

# Developing a GRMHD code for heterogeneous computing

## Challenges and perspectives



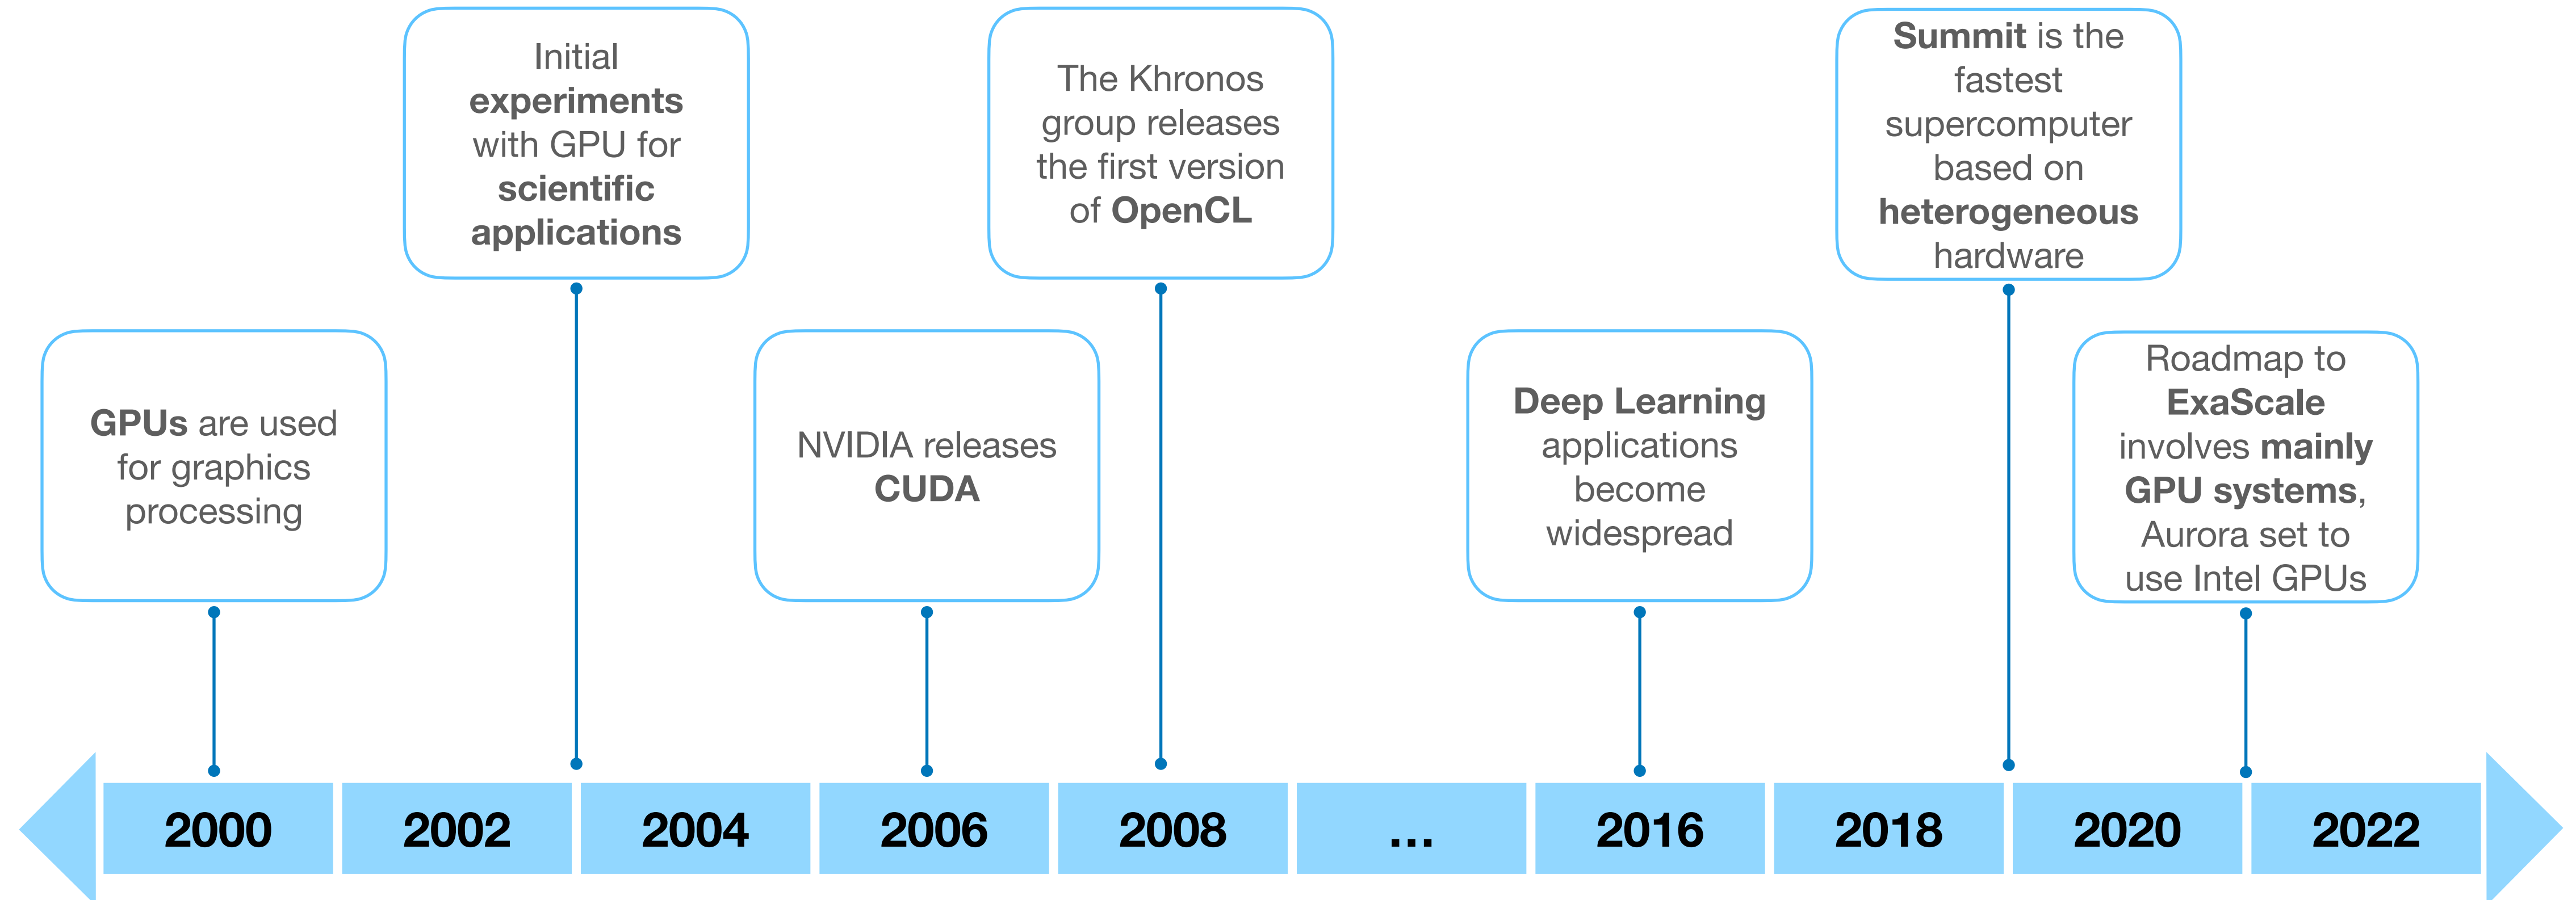
Based on work done in collaboration with L. Rezzolla

# Table of Contents

1. Introduction to GPU computing
2. Application to GRMHD equations

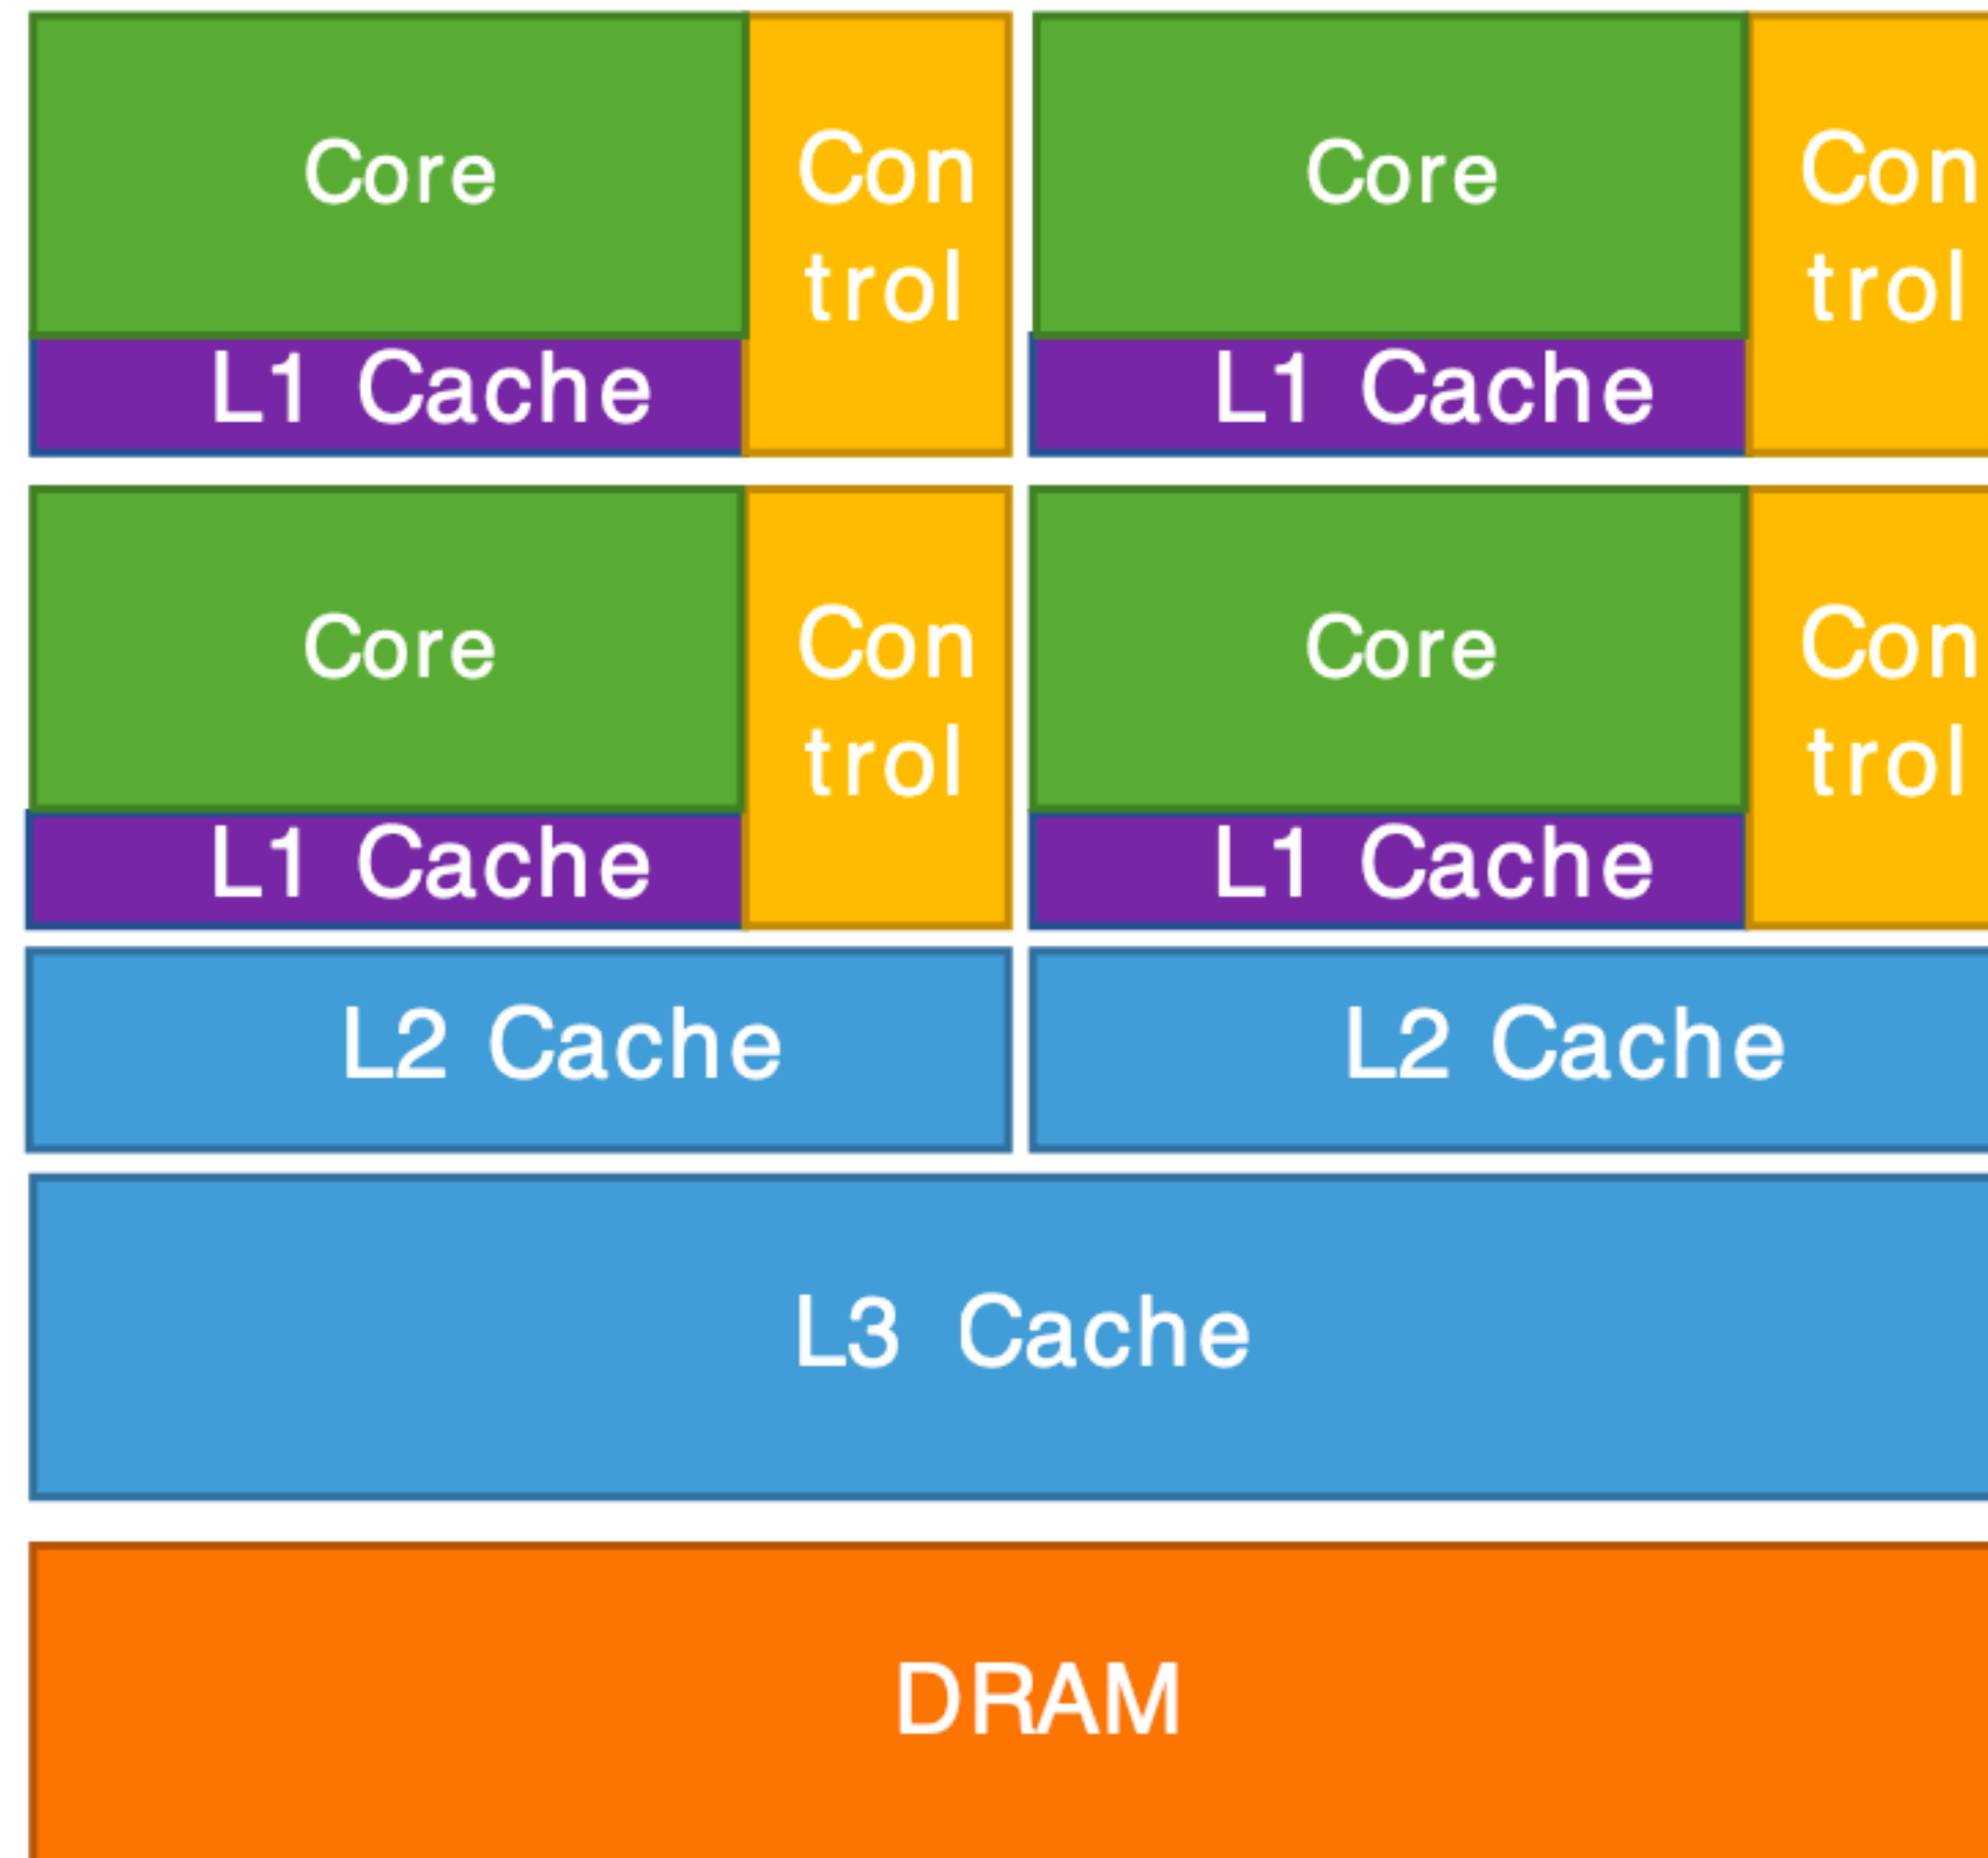
# Introduction and motivation

**Graphics Processing Units** are becoming prevalent tools for High Performance Computing




# Introduction and motivation

GPUs on paper offer **far more raw compute power** than CPUs



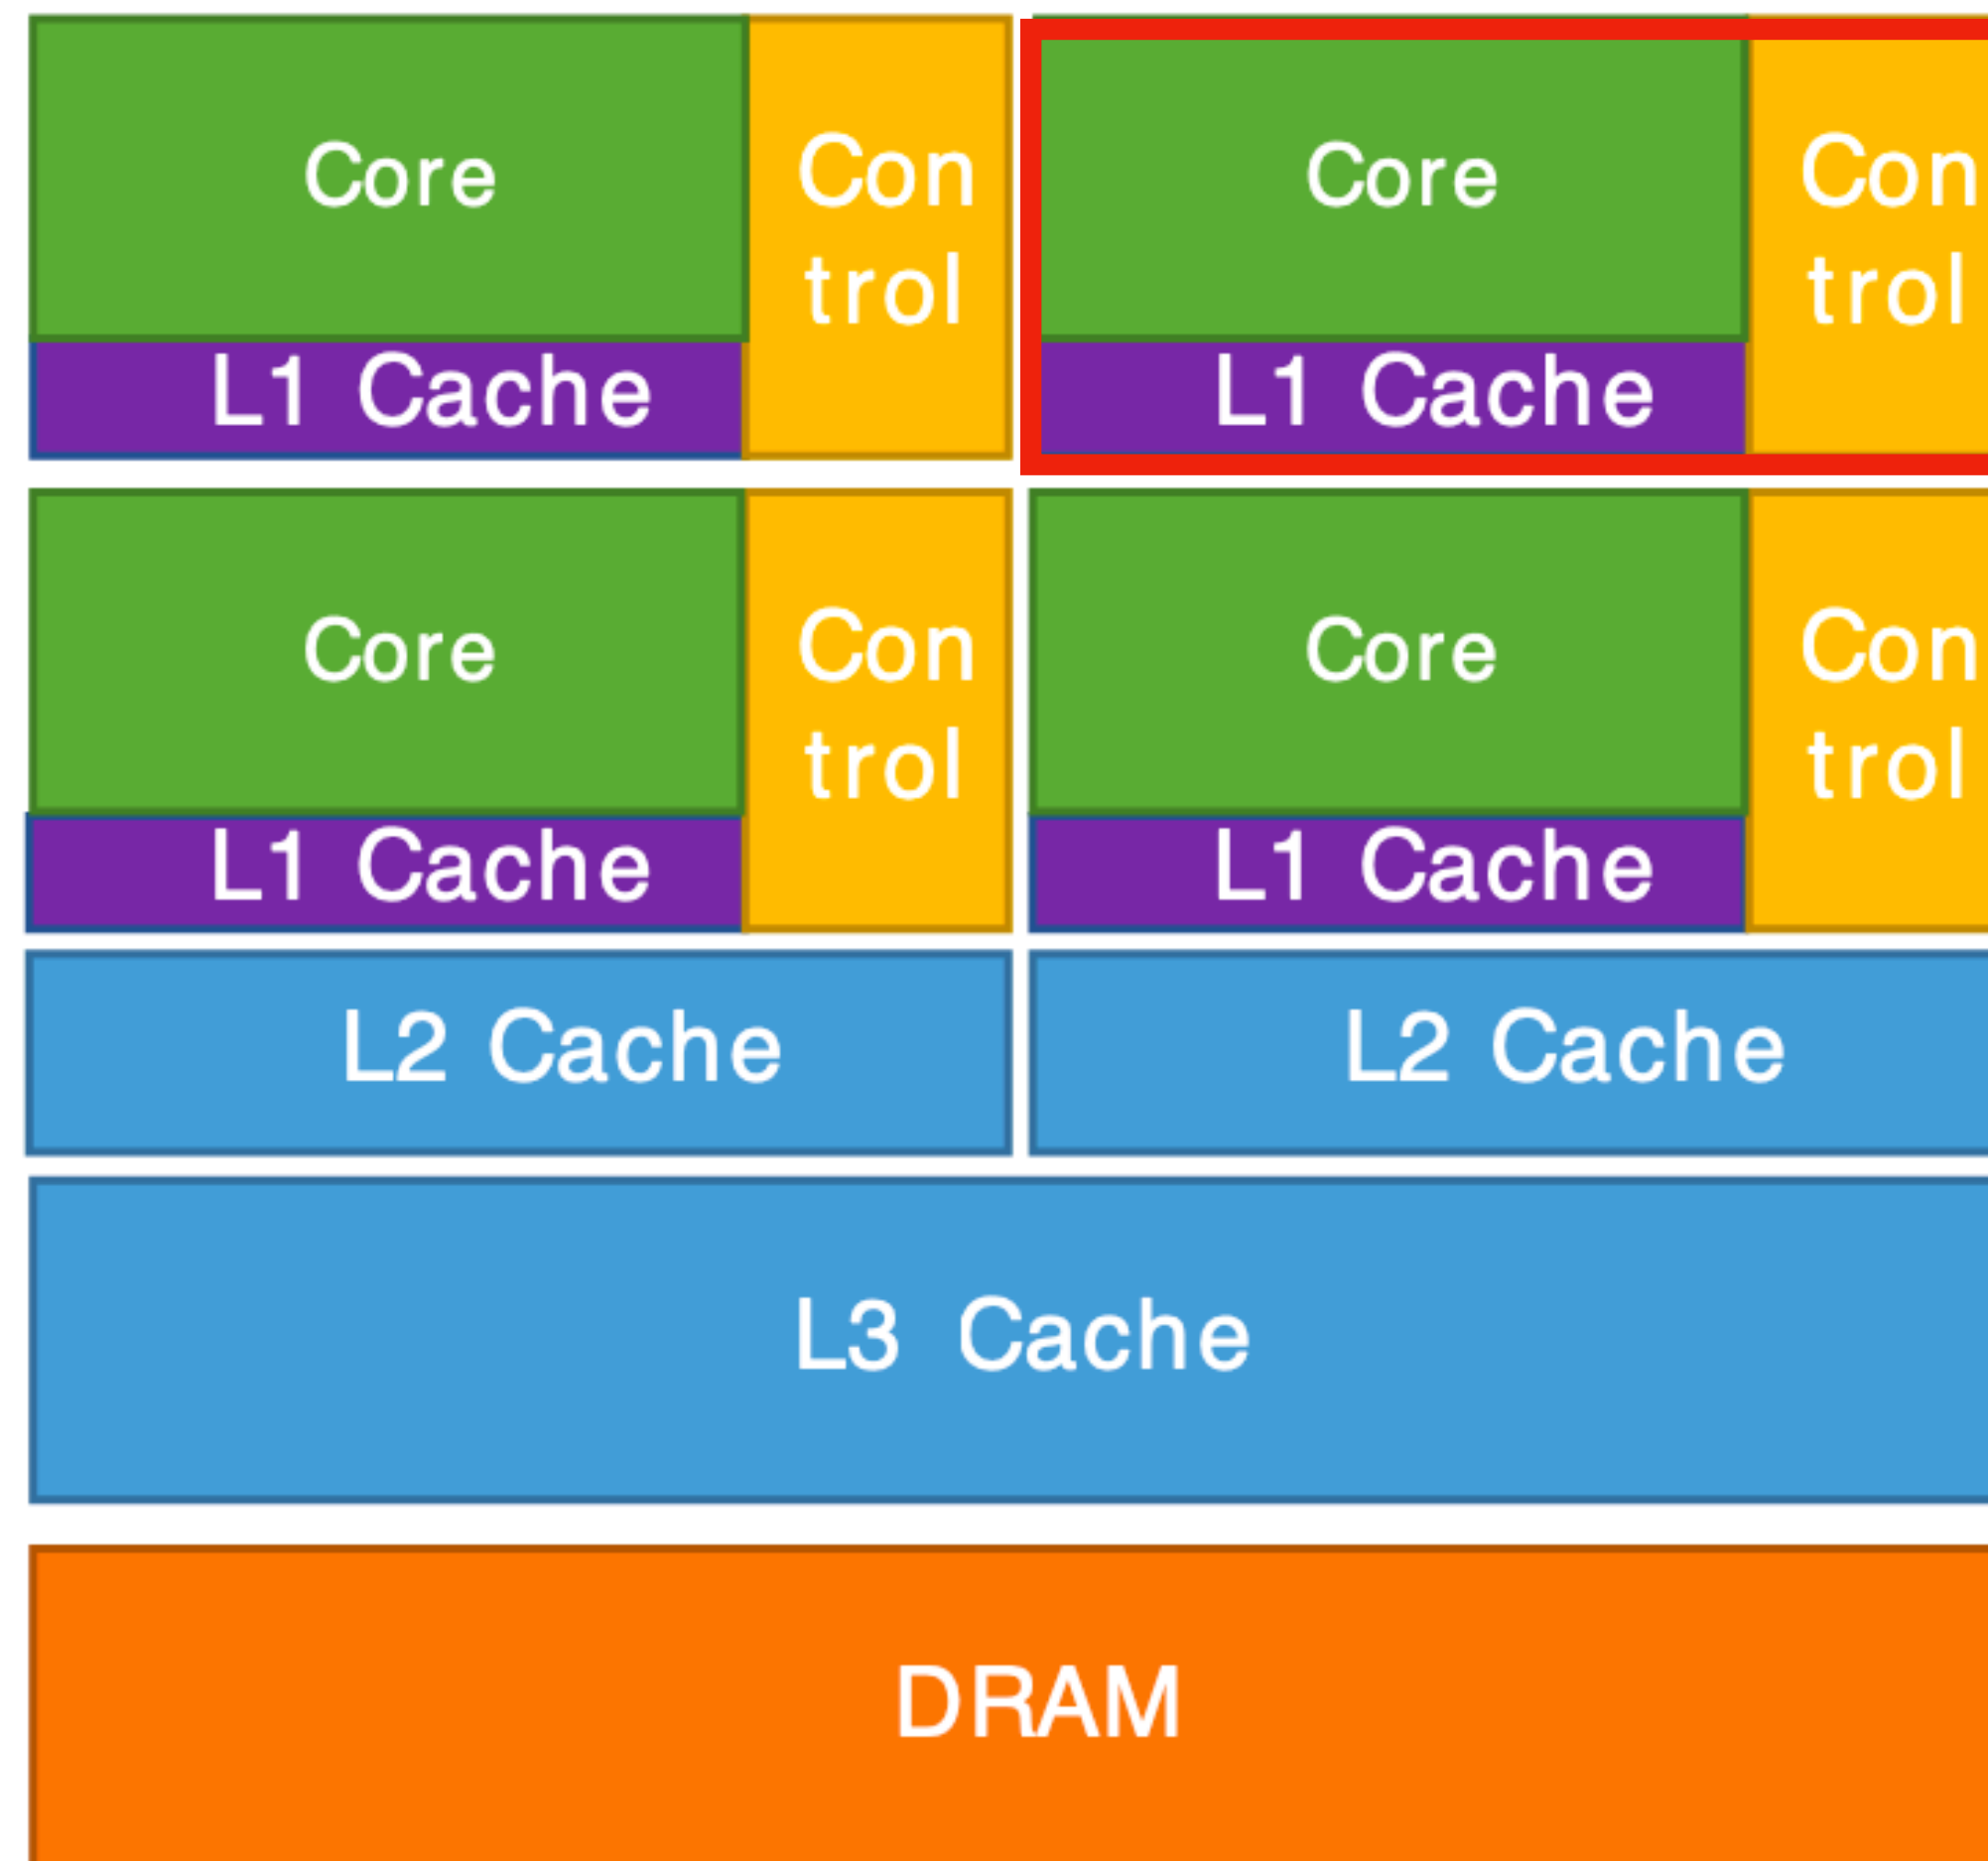
CPU



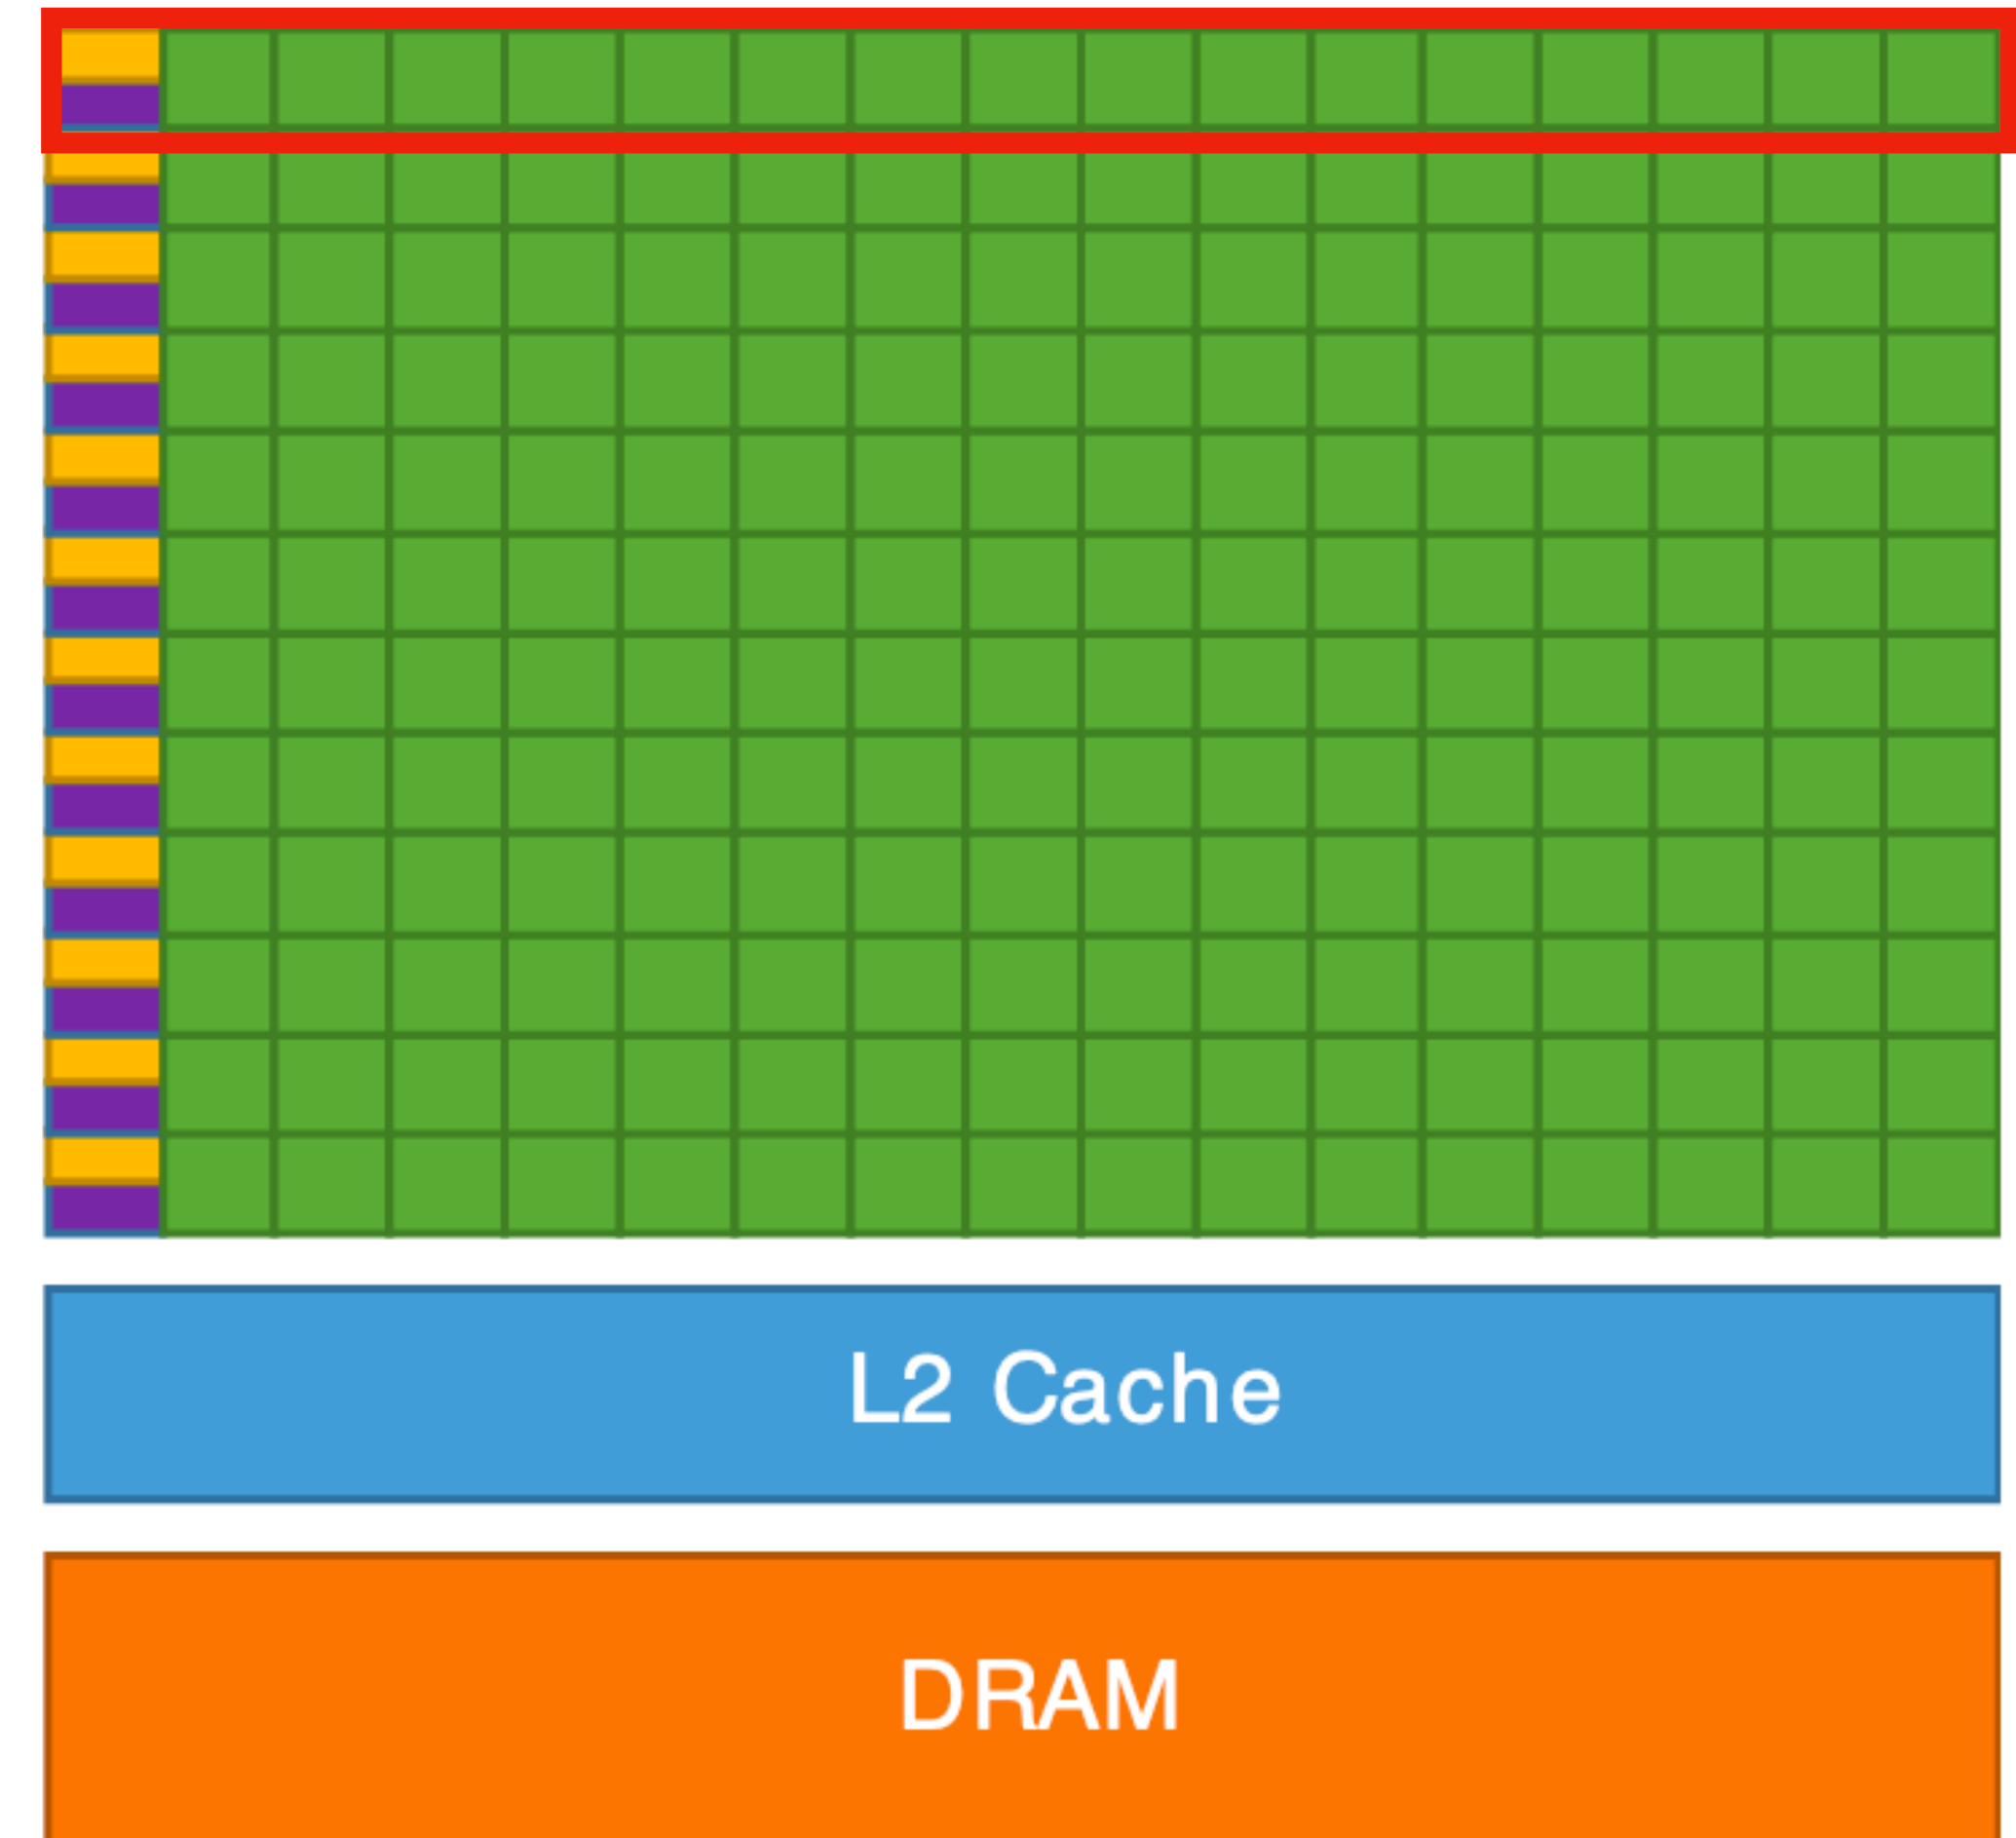
GPU

# Introduction and motivation

The microarchitecture and data models are **very different**.



CPU



GPU

# What works well on GPUs

- **Ray tracing** and image processing

High degree of **parallelism**, little data **dependency**

Repeat for N rays

Trace a single ray

- **AI and Machine Learning** (e.g. deep neural net back propagation)

Lots of operations can be run **independently**, mostly reliant on **linear algebra kernels**

Repeat for N neurons

Repeat for M weights

Compute weight gradients



# What about PDE solvers?

## CONS

- I) **Hyperbolic PDEs** describe the transport of information.
  - a. **Causal structure** —> **data dependencies**
  - b. Need for **communication** and **synchronization**
- II) Large **I/O** and **memory** requirements
- III) GPUs are usually **optimized** for FP8-16-32 workloads.

## PROS

- I) Plenty of **parallelism**, lots of **grid sites** / **particles** to update
- II) Can benefit heavily from **SMP** (~shared memory parallelism)
- III) Mixed **hyperbolic** / **elliptic** systems could have even **larger** benefits.

# What about PDE solvers?

Only **one way to know** for sure.

-> We are developing a **new GRMHD framework** on GPU backends aimed at:

- Exploring the **applicability** of heterogeneous computing to computational astrophysics.
- Building a **modern** and **future-proof** tool for research.

Codename: **General Relativistic Astrophysics Code for Exascale**.



# Rest of this talk

- 1. Introduction to GPU computing**
- 2. Application to GRMHD equations**
  - Introduction to GRACE**
  - Code Tests & Preliminary Results**
  - What can we say about performance?**

# Introduction to GRACE

Two main components:

1. **p4est** AMR library

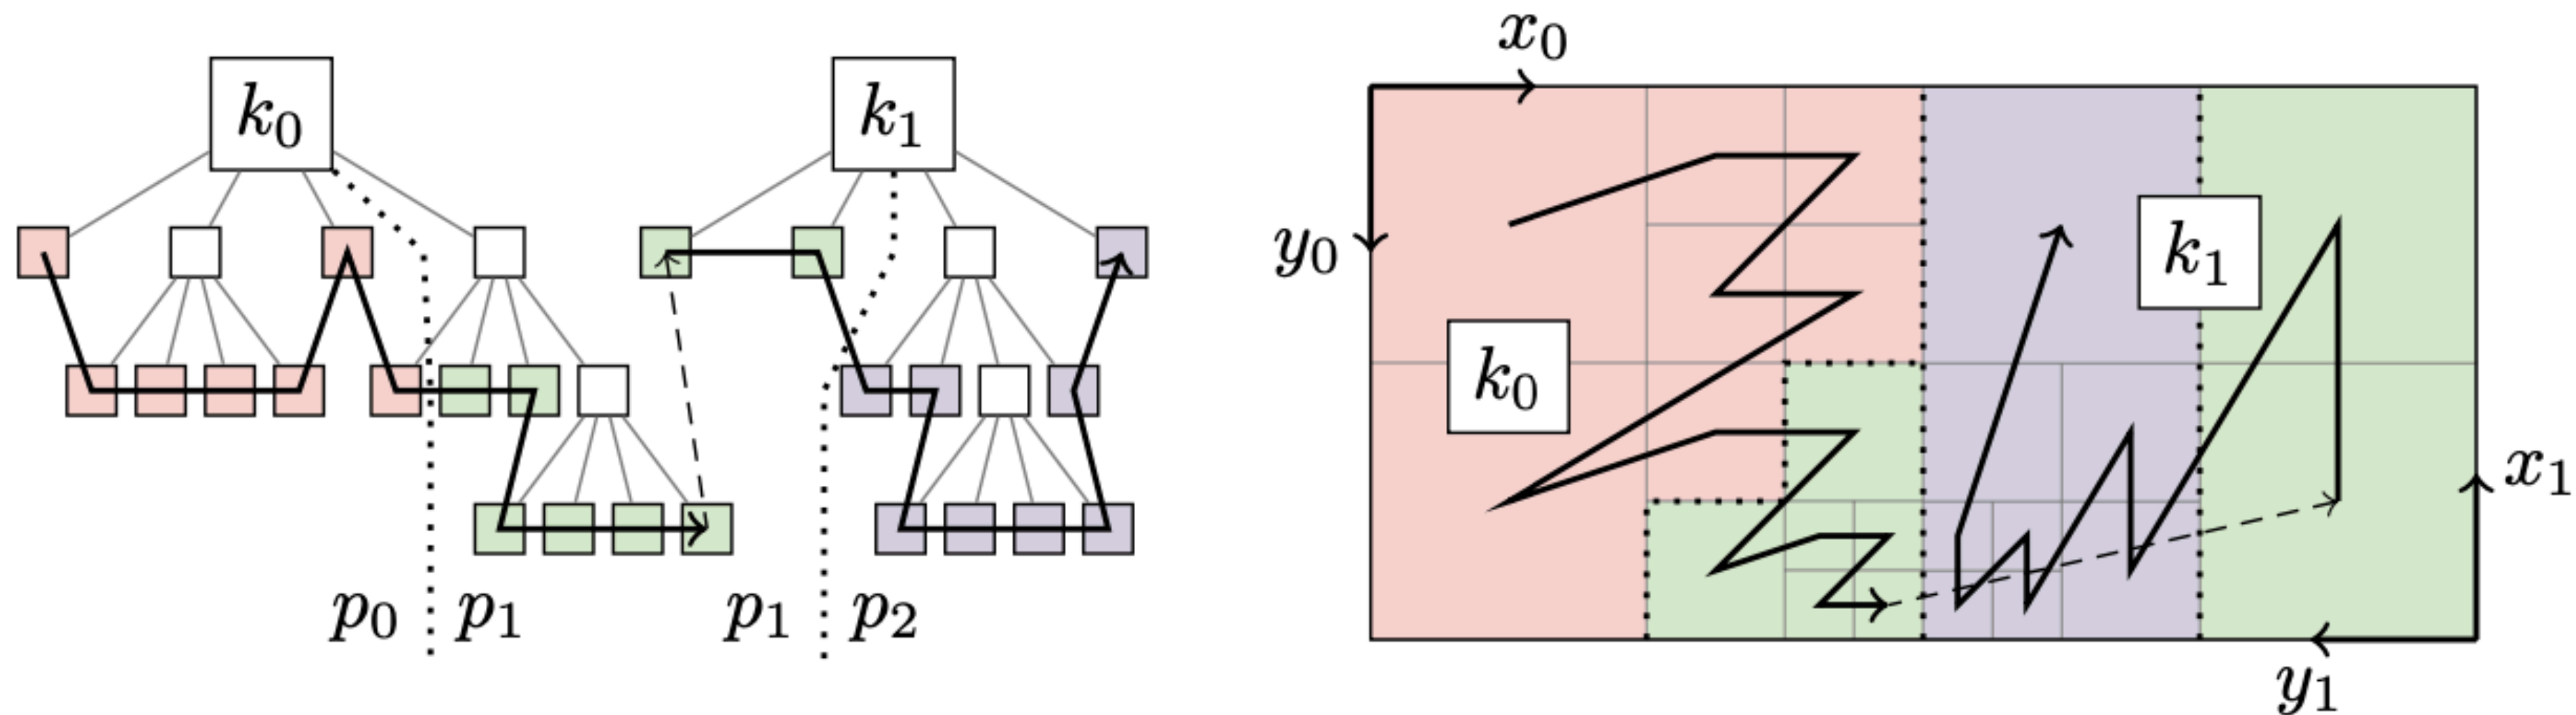
Grids with adaptive resolution are a fundamental ingredient of any code that aims at serious scientific contributions.

2. **Kokkos** Performance Portability Layer

GPUs are complex and varied (different vendors, different APIs) and a software layer in between the physics code and the silicon helps to mitigate these challenges.

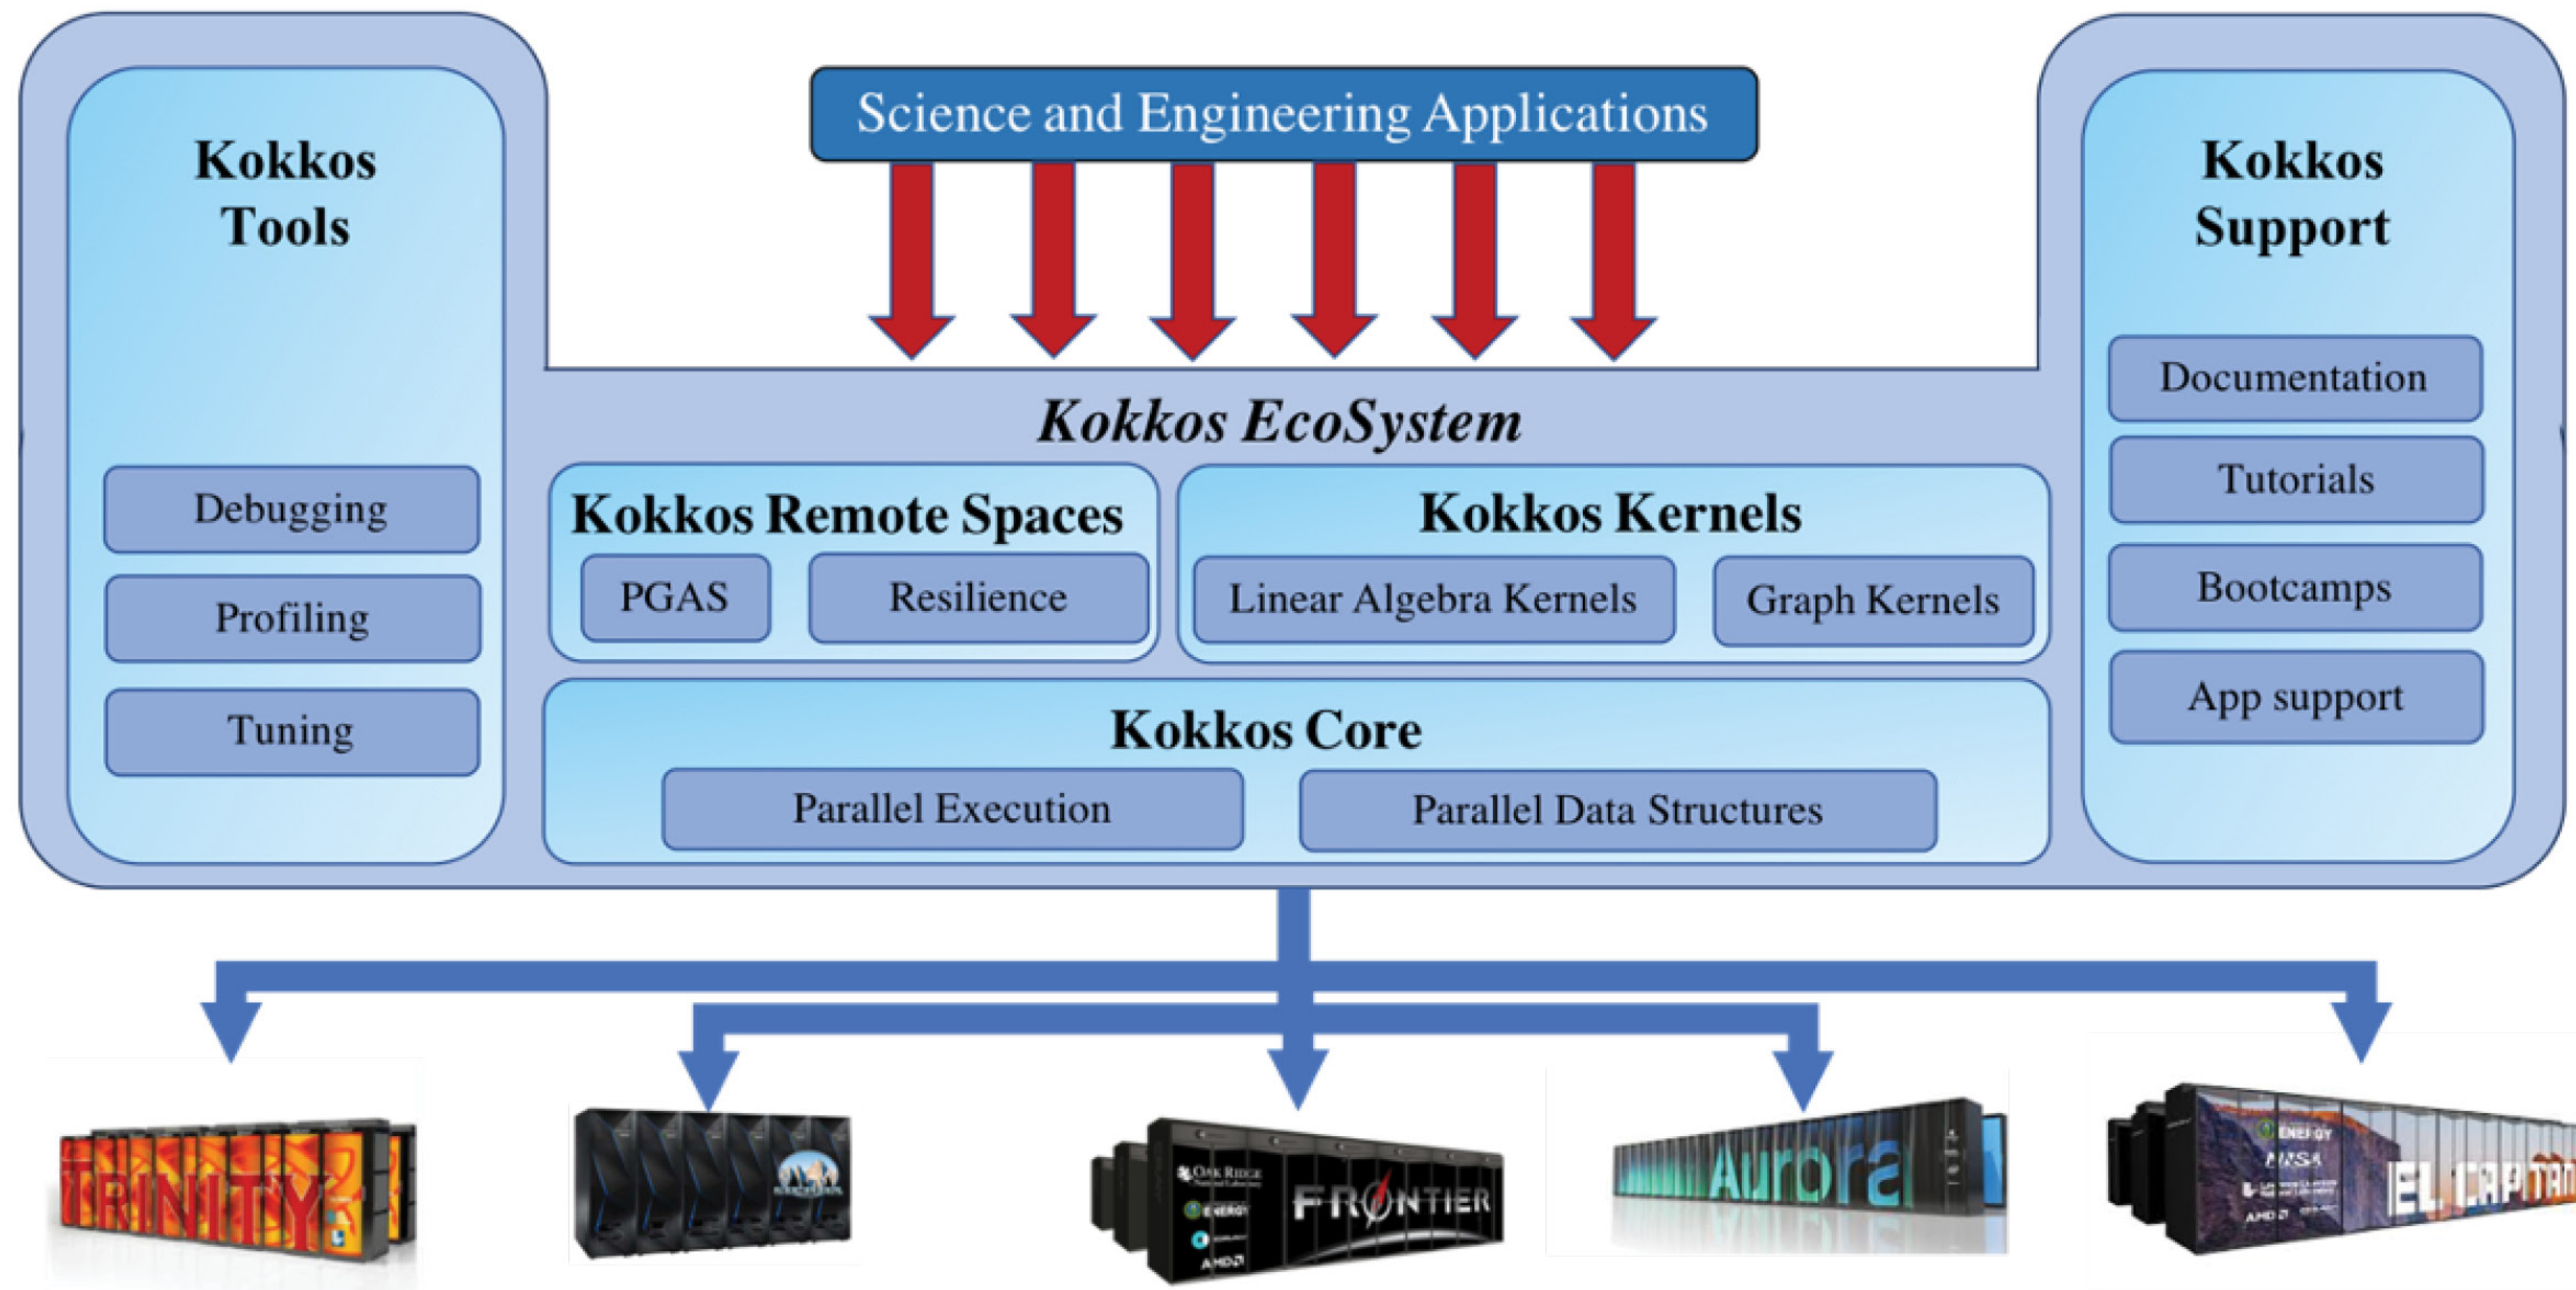
# p4est AMR library

- **p4est** only handles the grid, **not** the data
- More dev work but ideal for GPUs
- Data always sits on the GPU
- AMR routines (prolongation, restriction, ghost zones) are custom written **GPU kernels**



# Kokkos Performance Portability Layer

**Kokkos** is used in GRACE to **offload** work to GPU devices.





# Code Tests & Preliminary Results

Two **model** equations:

*i) **Scalar advection***

Simplest **hyperbolic** equation. Test of basic finite-volume + AMR infrastructure.

*ii) **Burgers equation***

**Nonlinear** hyperbolic PDE  $\rightarrow$  **High Resolution Shock Capturing** methods required to handle discontinuities.

**GRMHD** module currently under testing

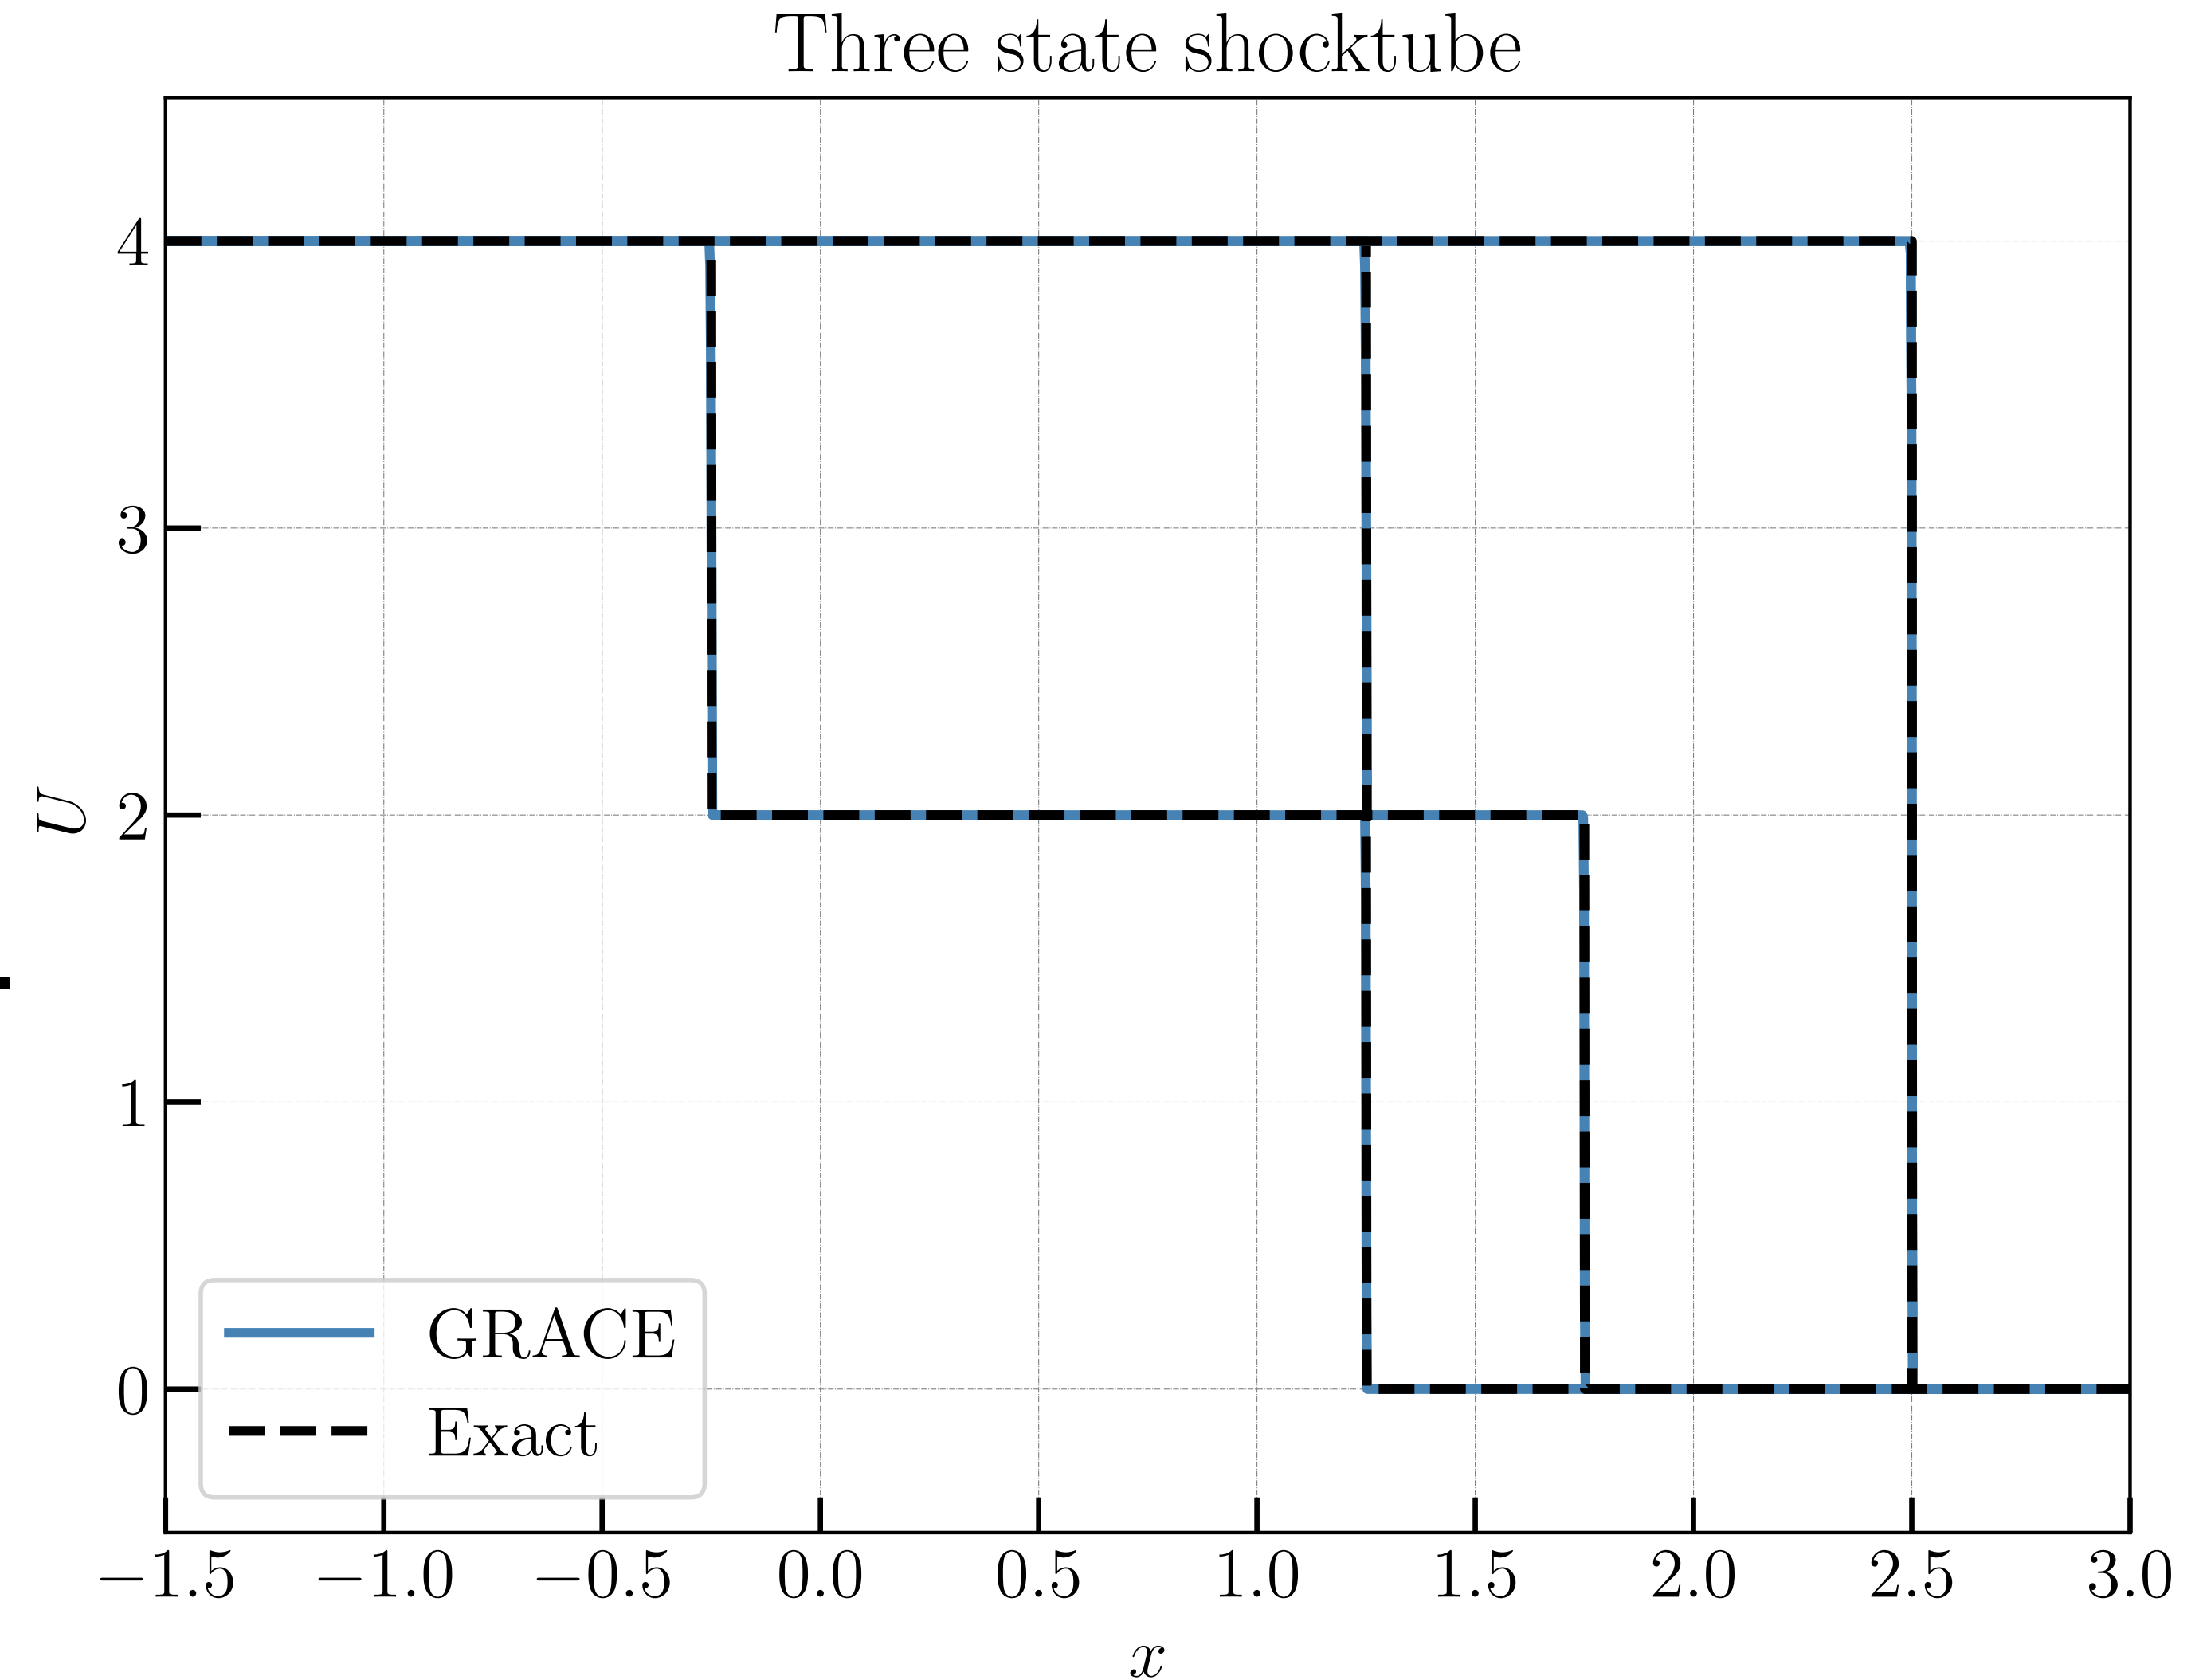
# Burgers equation

- **HRSC** solvers for **nonlinear** PDEs implemented on **Cartesian** grid (2D and 3D).
- Currently supported reconstruction algorithms: **minmod**, **monotonized-central**, **WENO (3rd/5th order)**.
- Currently supported Riemann solvers: **HLLE**.
- Prototypical PDE system: **Burgers' inviscid** equation.

$$\partial_t U(\mathbf{x}, t) + \frac{1}{2} \partial_x U(\mathbf{x}, t)^2 = 0$$

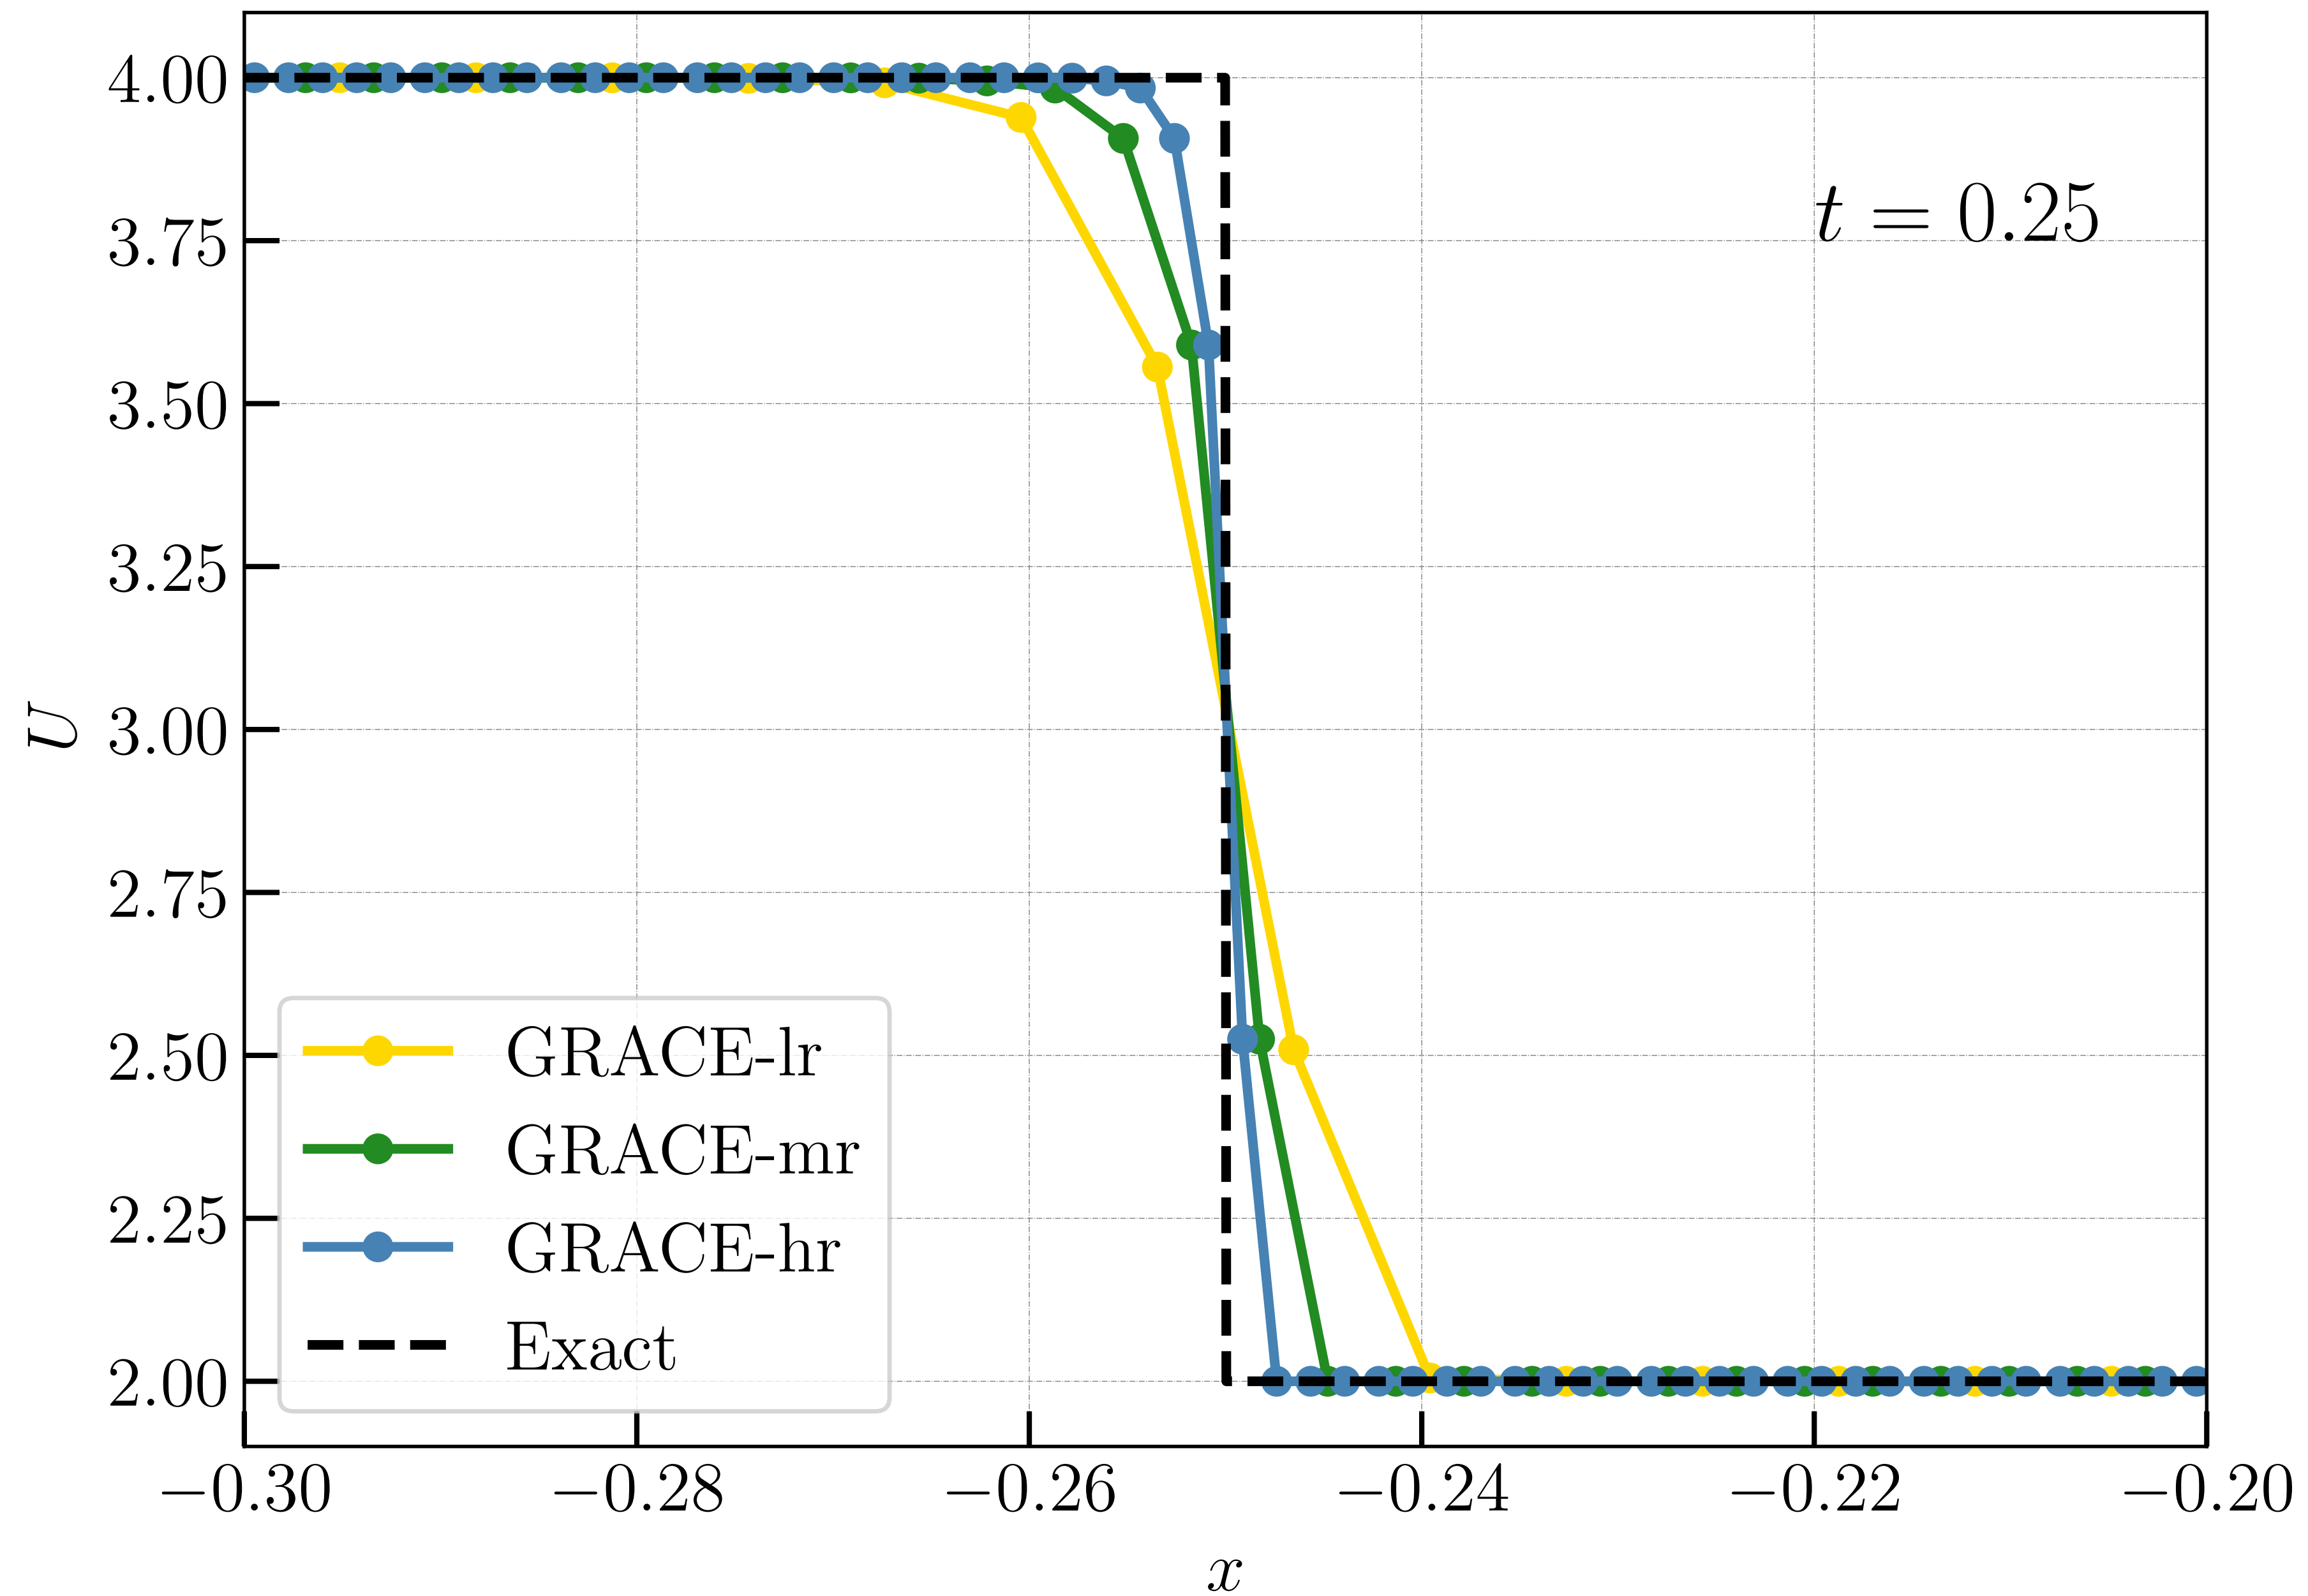
# Code Tests & Preliminary Results: Burgers equation

- Three-state **shock-tube** for Burgers' equation
- Solved in 3D with **uniform mesh** refinement and Runge-Kutta 2 time-stepping.
- The reconstruction method is **MC2**.
- This initial data leads to a double shockwave.





# Three state shocktube




# Code Tests & Preliminary Results

- **N-wave test**
- **Initial data**


$$U(x, 0) = e^{-\frac{(x-1)^2}{2}} - e^{-\frac{(x+1)^2}{2}}$$

- **Solution:**


$$U(x, t) = e^{-\frac{(x - U(x, t)t - 1)^2}{2}} - e^{-\frac{(x - U(x, t)t + 1)^2}{2}}$$



# Results



# Order of Convergence



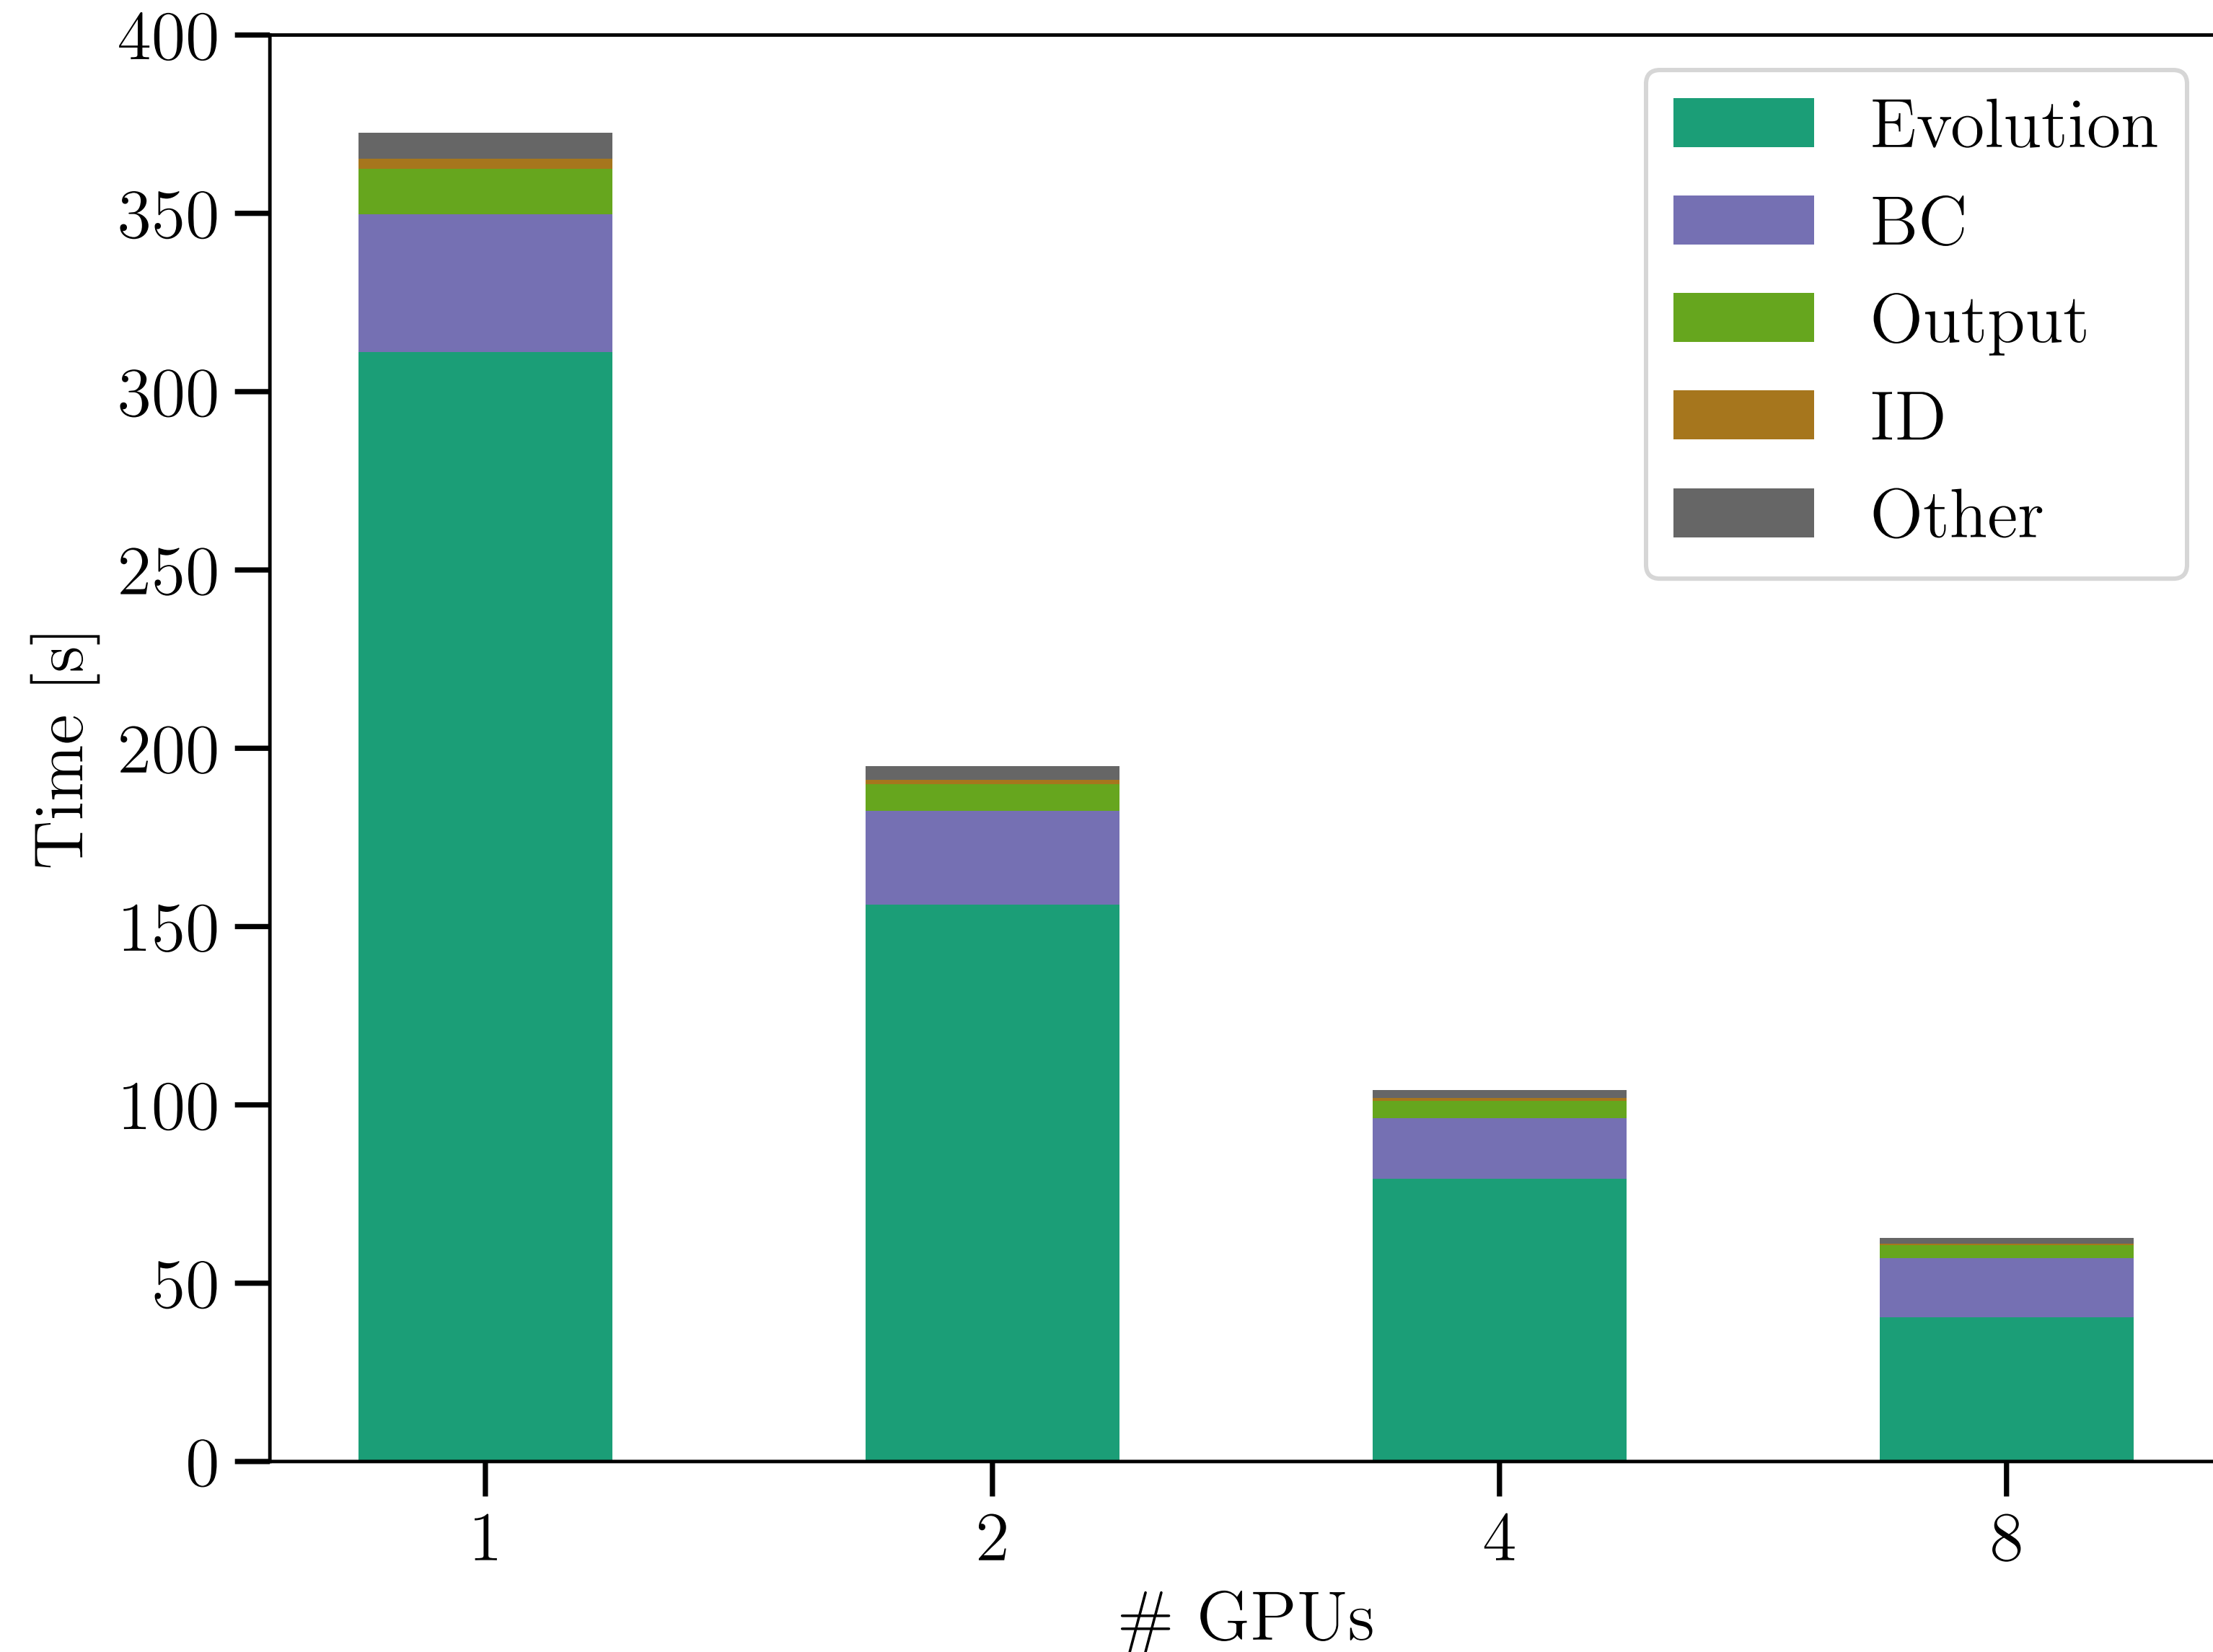
# What can we say about performance?

**Profiling** codes on heterogeneous systems is **hard**.

- Scaling is a measure of how well an application performs on a **large number** of compute resources
  - **Strong scaling**: fixed problem size, increasing resources
  - **Weak scaling**: problem size grows proportionally to the available resources
  - Scaling alone is not always a good proxy for performance.
- Code **efficiency** on a single compute unit is largely **uncorrelated** to **scaling**
  - Measuring efficiency can be very **challenging**
  - Having a grip on “single-core” code performance is key for **effective optimization**.

# Profiling case study: Unigrid simulation



- Unigrid simulation of Burgers equation N wave test case.
- 5 levels of refinement with  $16 \times 16 \times 16$  points / block + ghost zones.
- **Take-away:** Over 80% of the time is spent doing useful calculations





# Profiling case study: Unigrid simulation

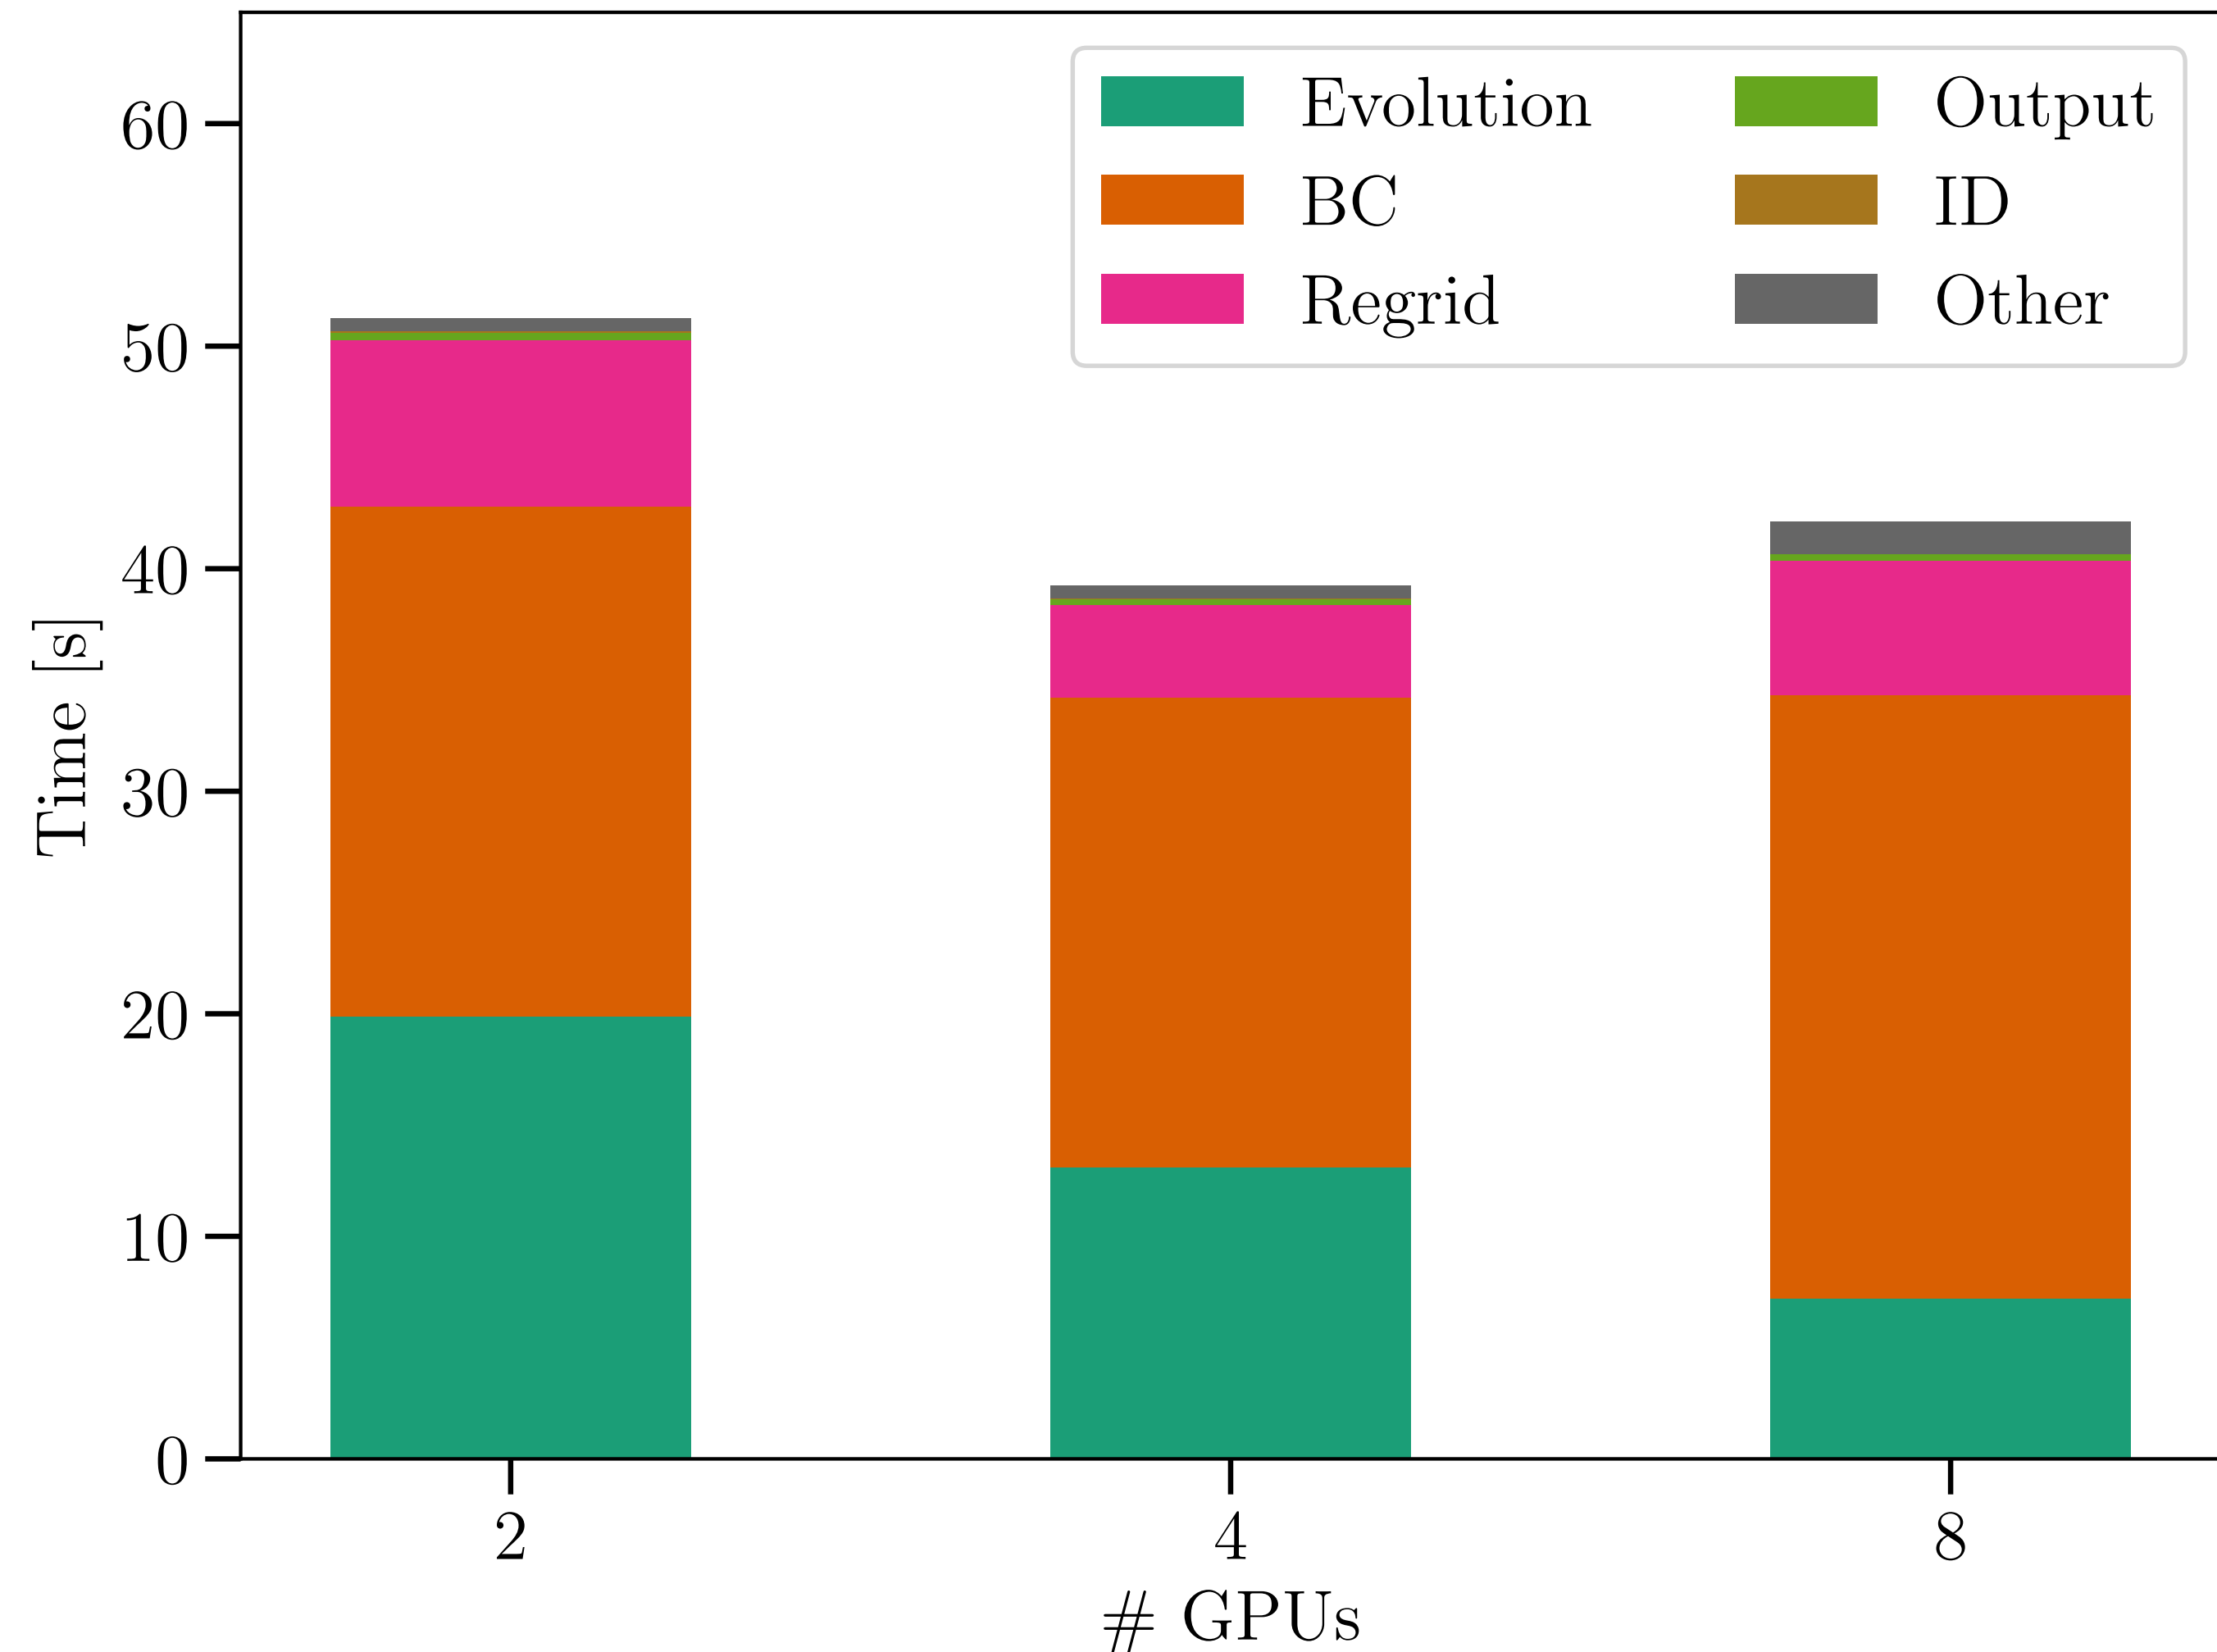
- The code scales well on the (limited) available resources.
- **Caveat:** MPI not properly fine-tuned for the system.
- **Take-away:** need a production environment to properly assess performance.





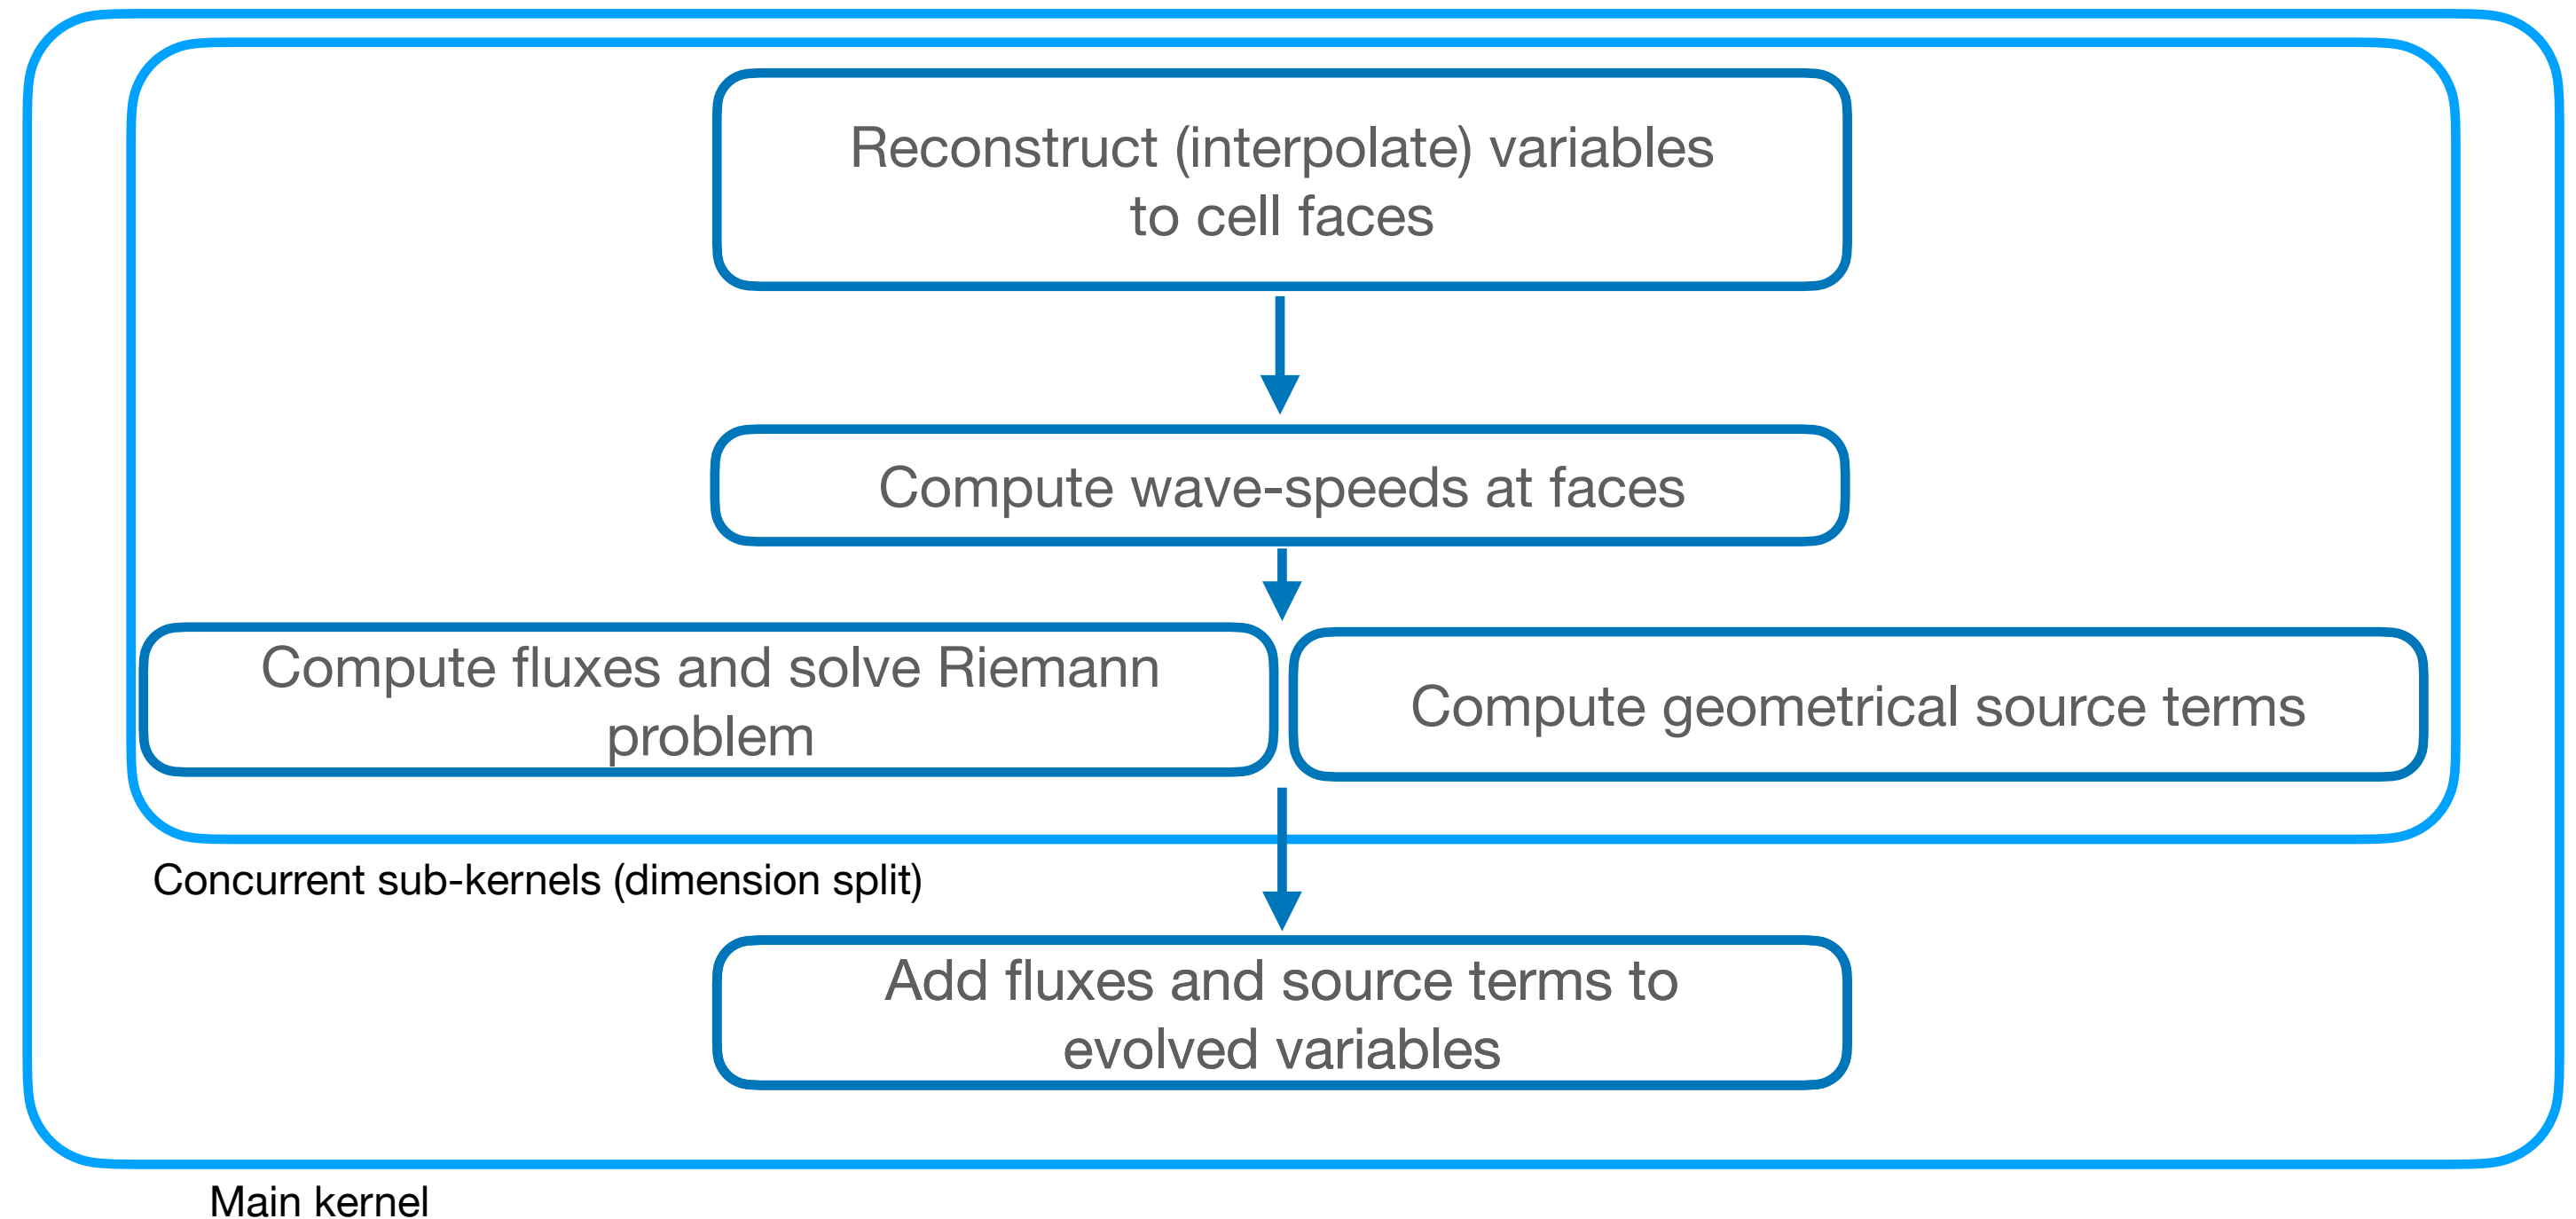
# Profiling case study: AMR simulation

- AMR simulations of Gaussian pulse advection with periodic boundaries
- 5 levels of refinement with  $16 \times 16 \times 16$  points / block + ghost zones.
- **Take-away:**  
Hanging interfaces are costly to handle



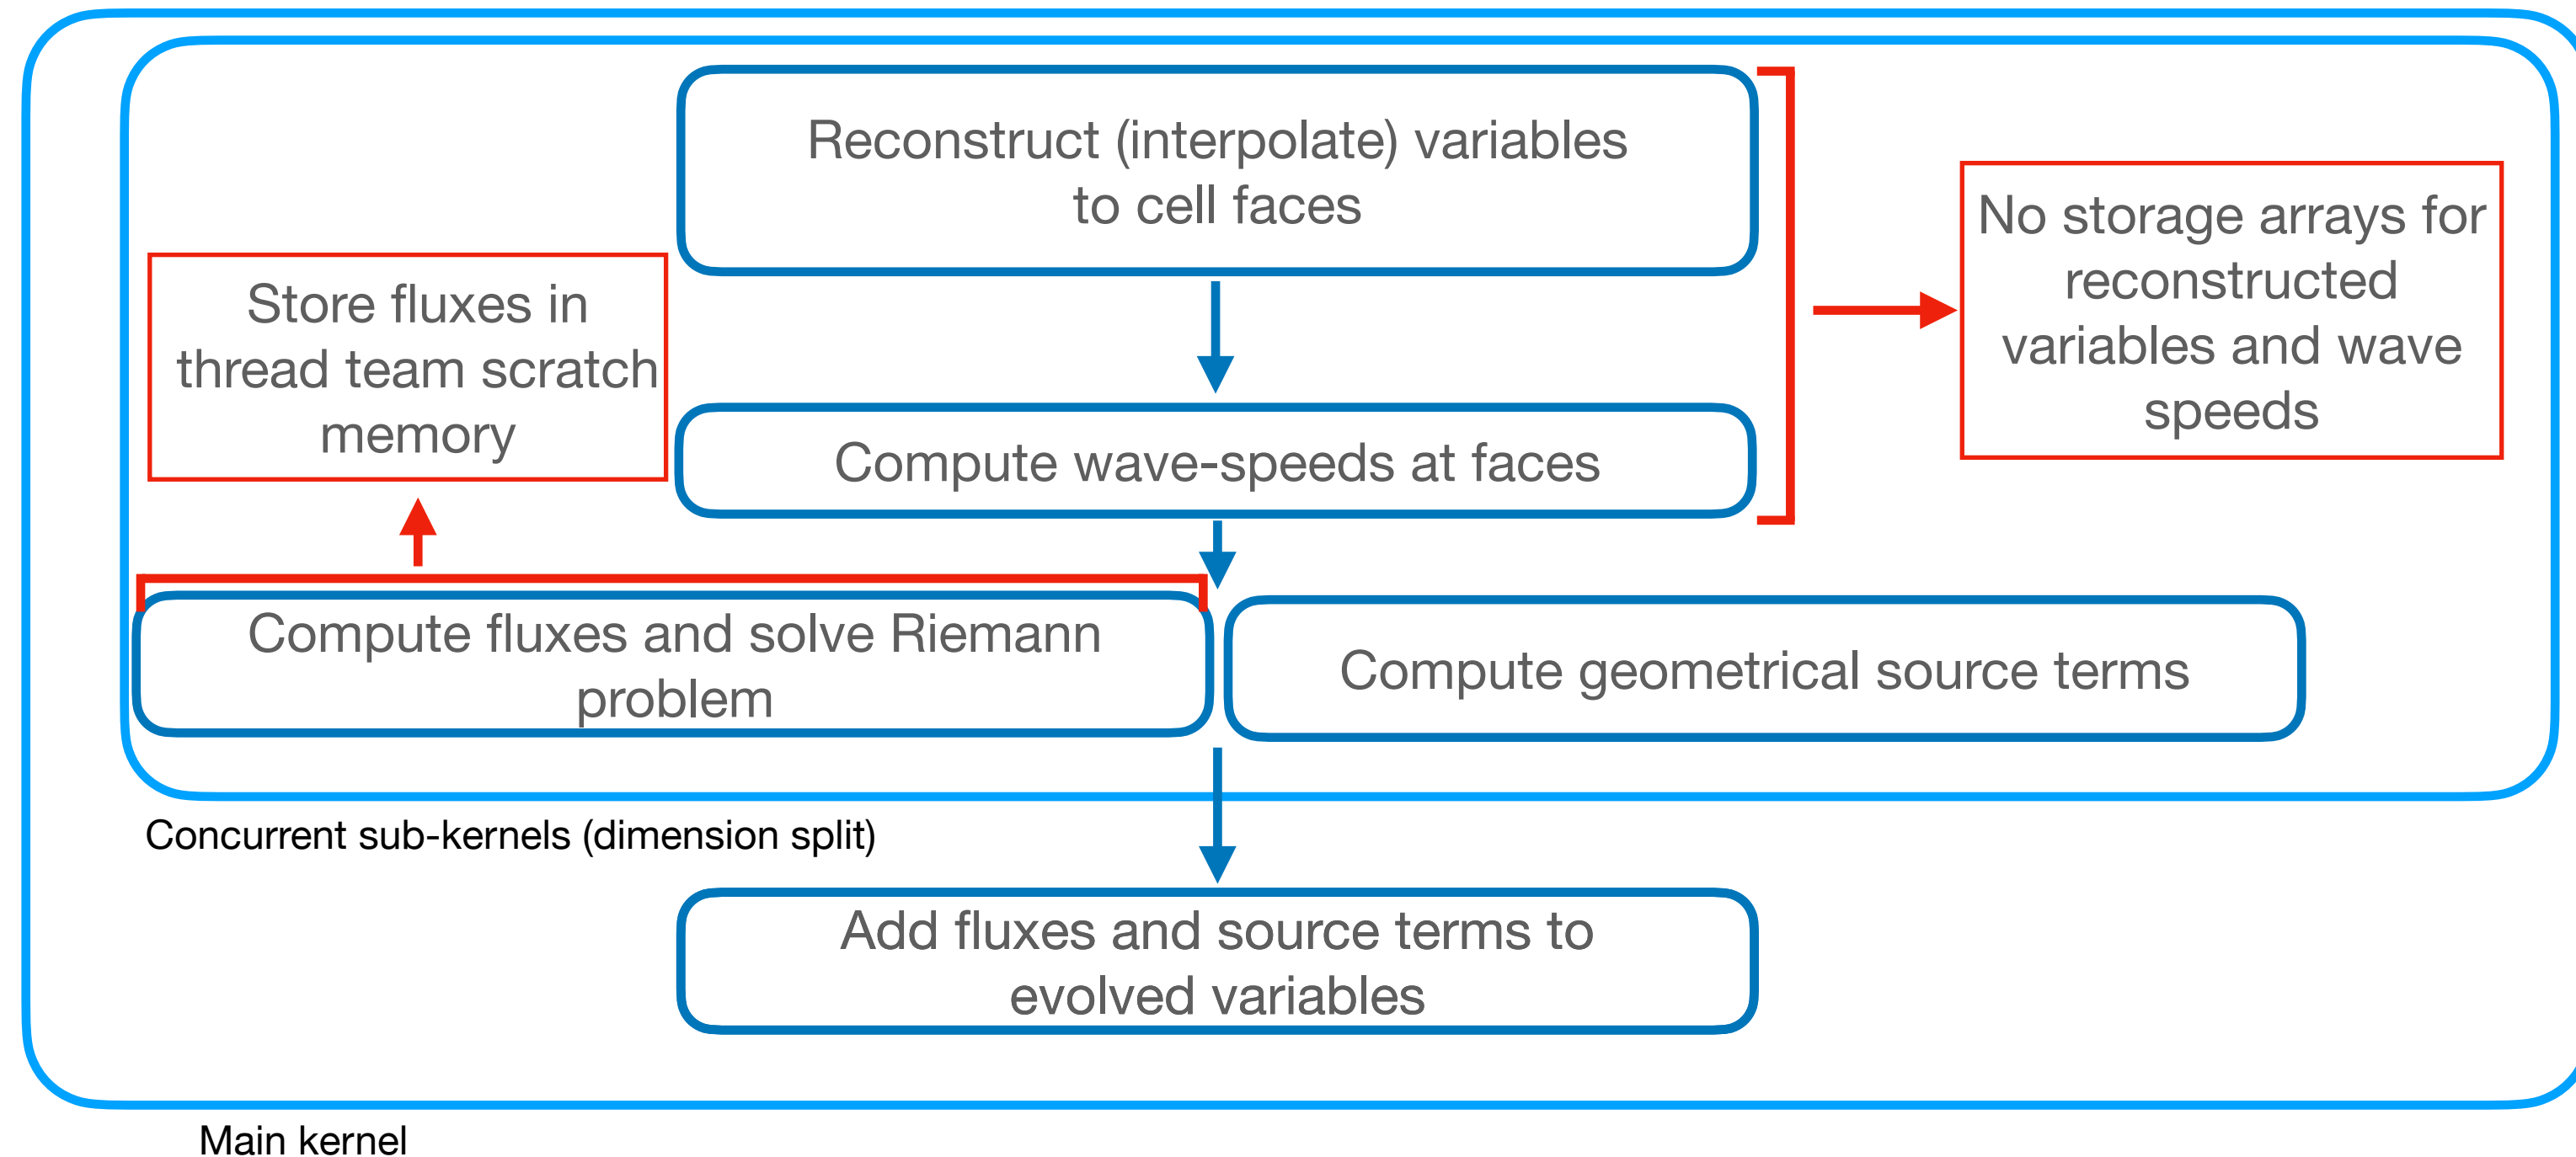
# Code performance: evolution kernel

- Evolution: most **time-intensive** section.
- **Schematically** consists of a series of directional loops and a final loop to add sources
- Performance counters sampled with low-level device profilers



# Code performance: evolution kernel


- **Heuristic** optimization based on **saving memory** transfers where possible
- **Remove** intermediate **storage** at the price of **extra computations**
- **Directional dependence** reduces vectorization performance



# Code performance: roofline model

- **Optimistic** “speed of light” model for resource utilization.
- Applies to a **single computational kernel**.
- Bottleneck either:
  - **Execution of work**  $\rightarrow P_{\text{peak}}$  [flop/s]
  - **Data path**  $\rightarrow I \cdot b_S$  [flop/byte x byte/s]

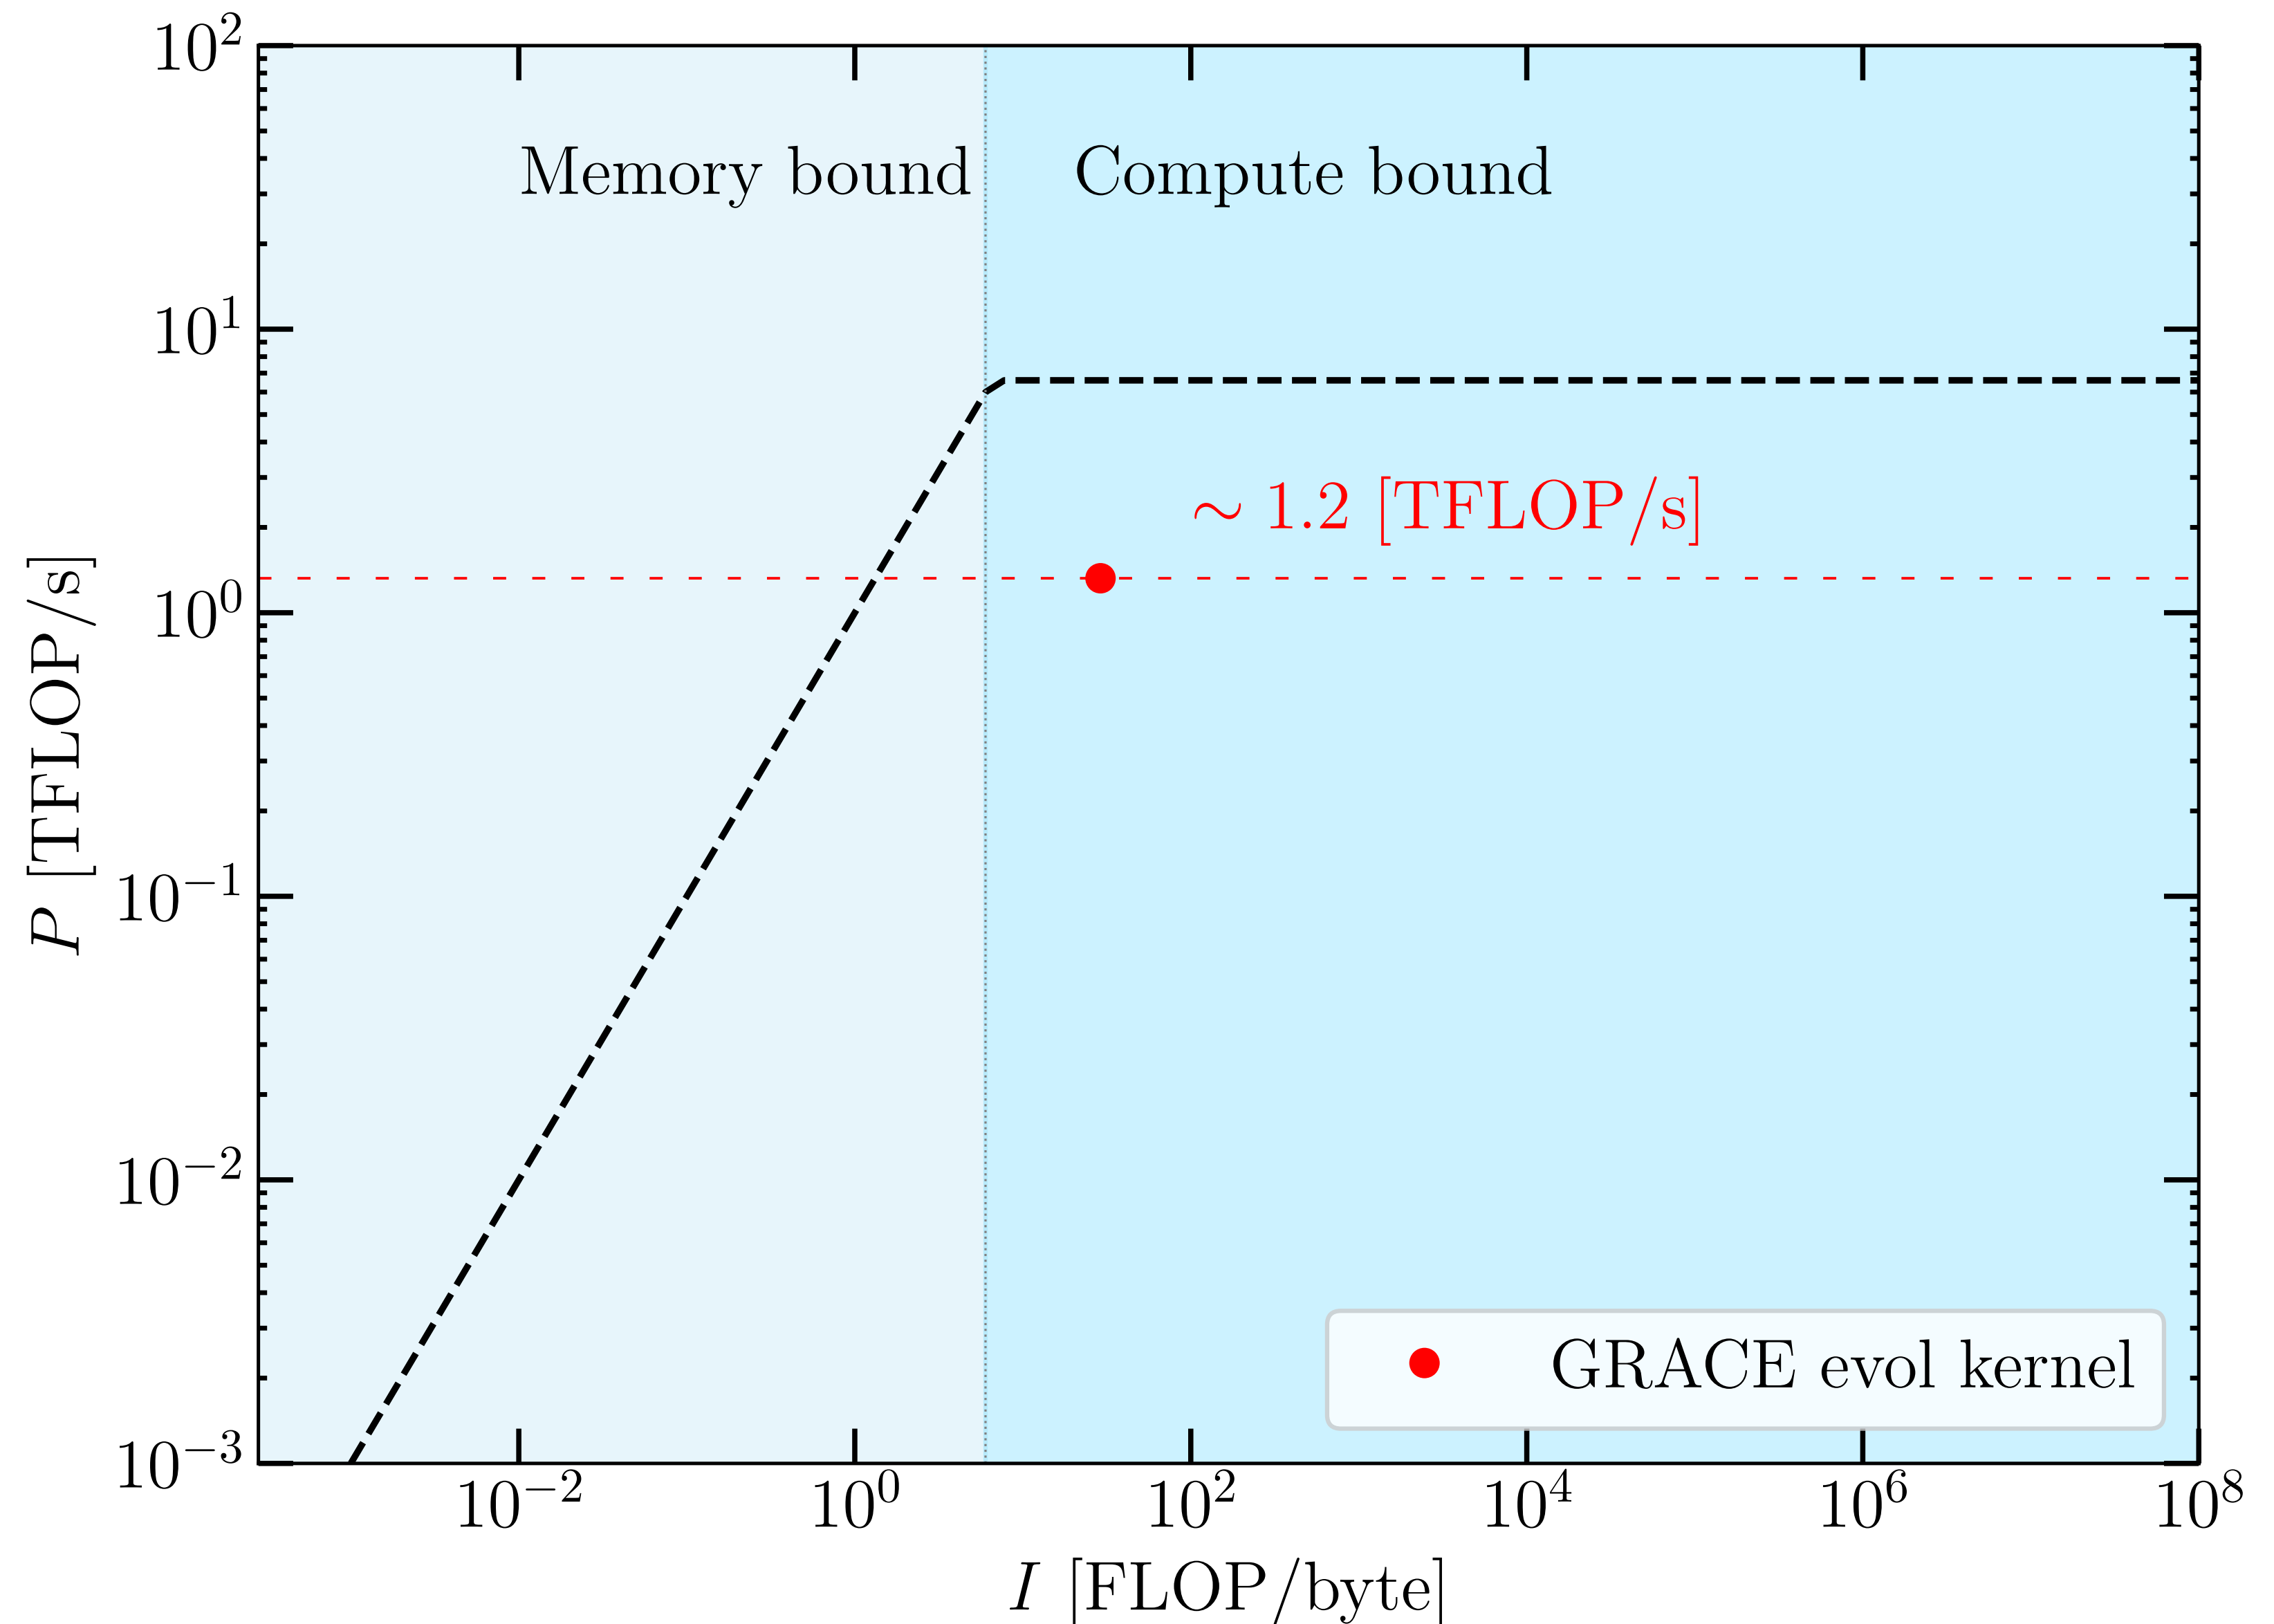
$$P = \min(P_{\text{peak}}, I \cdot b_S)$$



# Code performance: roofline model

- Evolution kernel **peaks** at  
~ **1.2 TFLOP/s**
- Tested on a single AMD  
Mi50 card
- Theoretical peak for **FP64**  
workload:  
  
~ **6.5 TFLOP/s**

$$P = \min(P_{\text{peak}}, I \cdot b_S)$$

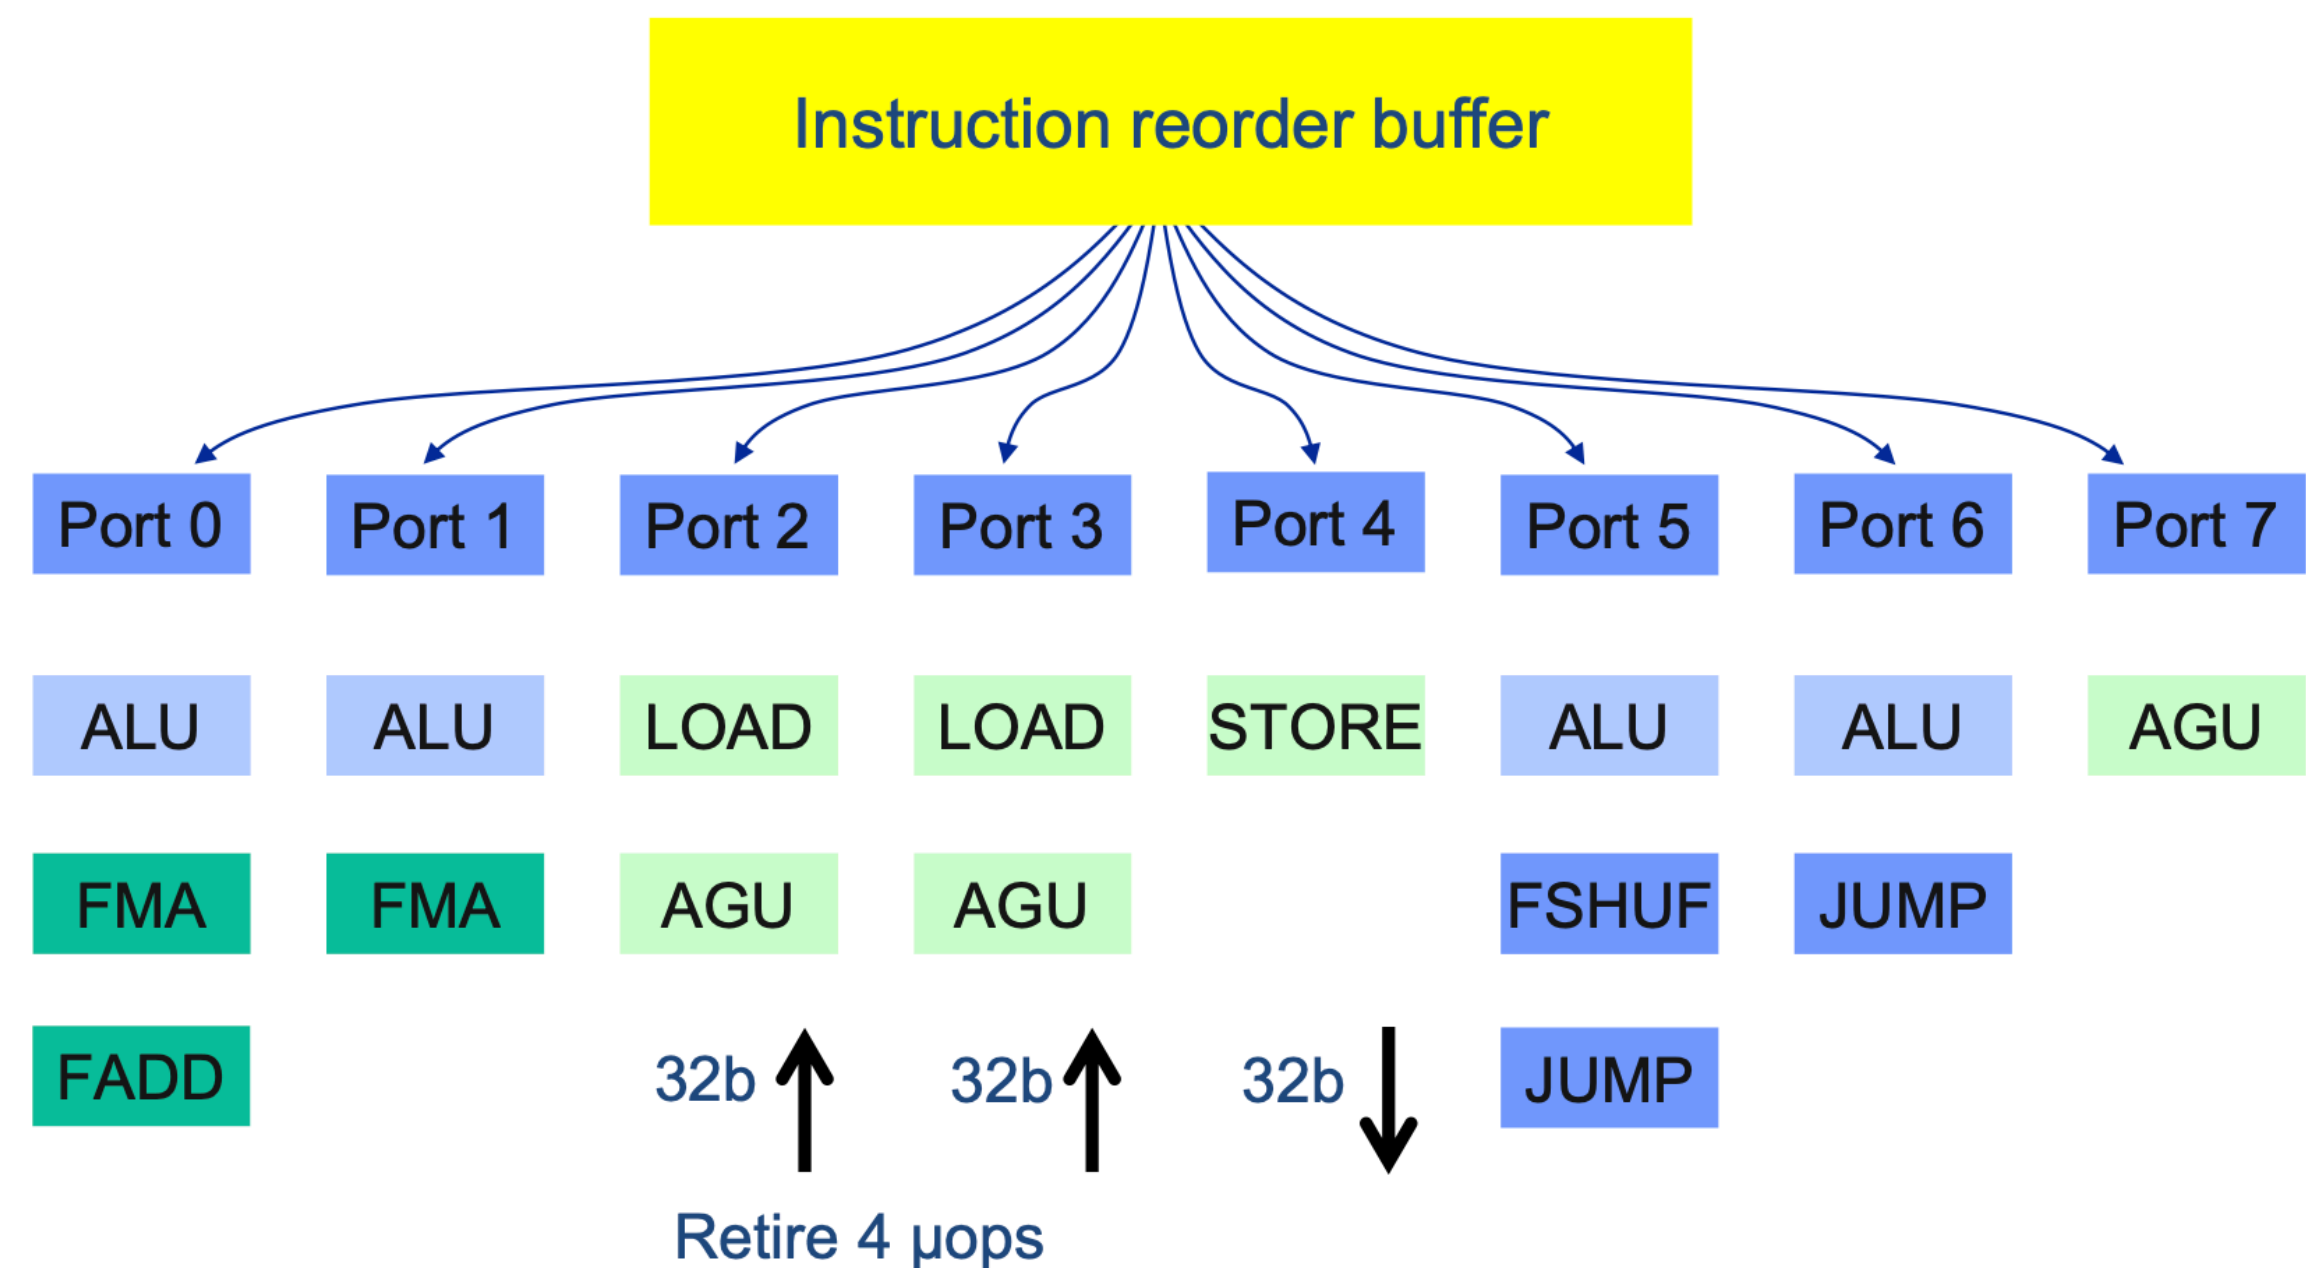





# Code performance: roofline model

Pinning down compute bottlenecks is **complex**.

- i. L2 **cache misses not** accounted for in simple model.
- ii. Likely suffering from **port contention**
- iii. **Vectorization** might be limited by loop dependencies



Model port scheduler  
diagram for Intel Haswell  
chip



Example of more **realistic** roofline  
model diagram

# Conclusions

- We have a **functional**, oct-tree **AMR enabled framework** for solving non-linear hyperbolic PDEs on GPU backends.
- The code shows **promising single-GPU performance** and **good scaling** on the available in-situ resources.
- The tool is fully built in-house and can be **easily extended**.
- Further developments are in progress and we hope to be production-ready in the Fall.



- Theoretical Astrophysics relies on **HPC** for realistic and accurate simulations of complex systems.
- Overwhelming majority of existing codes are **complex** and built upon years **iterative improvements / modifications**.
- The basic algorithms can often be **streamlined** by careful reviewing of existing codebases.
- The basic computing **paradigm** has **shifted** since these codes were designed.
- **GPUs** are an (almost) obvious candidate for modern simulation codes infrastructures.

- GPU offloading is **not** a drop-in replacement for traditional parallel programming paradigms.
- Current grid-based codes typically perform **very poorly** with a large number of **threads** (why?).
- Efficiency and performance require **knowledge** of both the basic **hardware architecture** and the **algorithm**.
- **Low-level** routines likely need to be revised to better suit the execution model of target machines.
- **High-level** algorithms also require modifications achieve the best possible performance on modern systems.

# Introducing GRACE

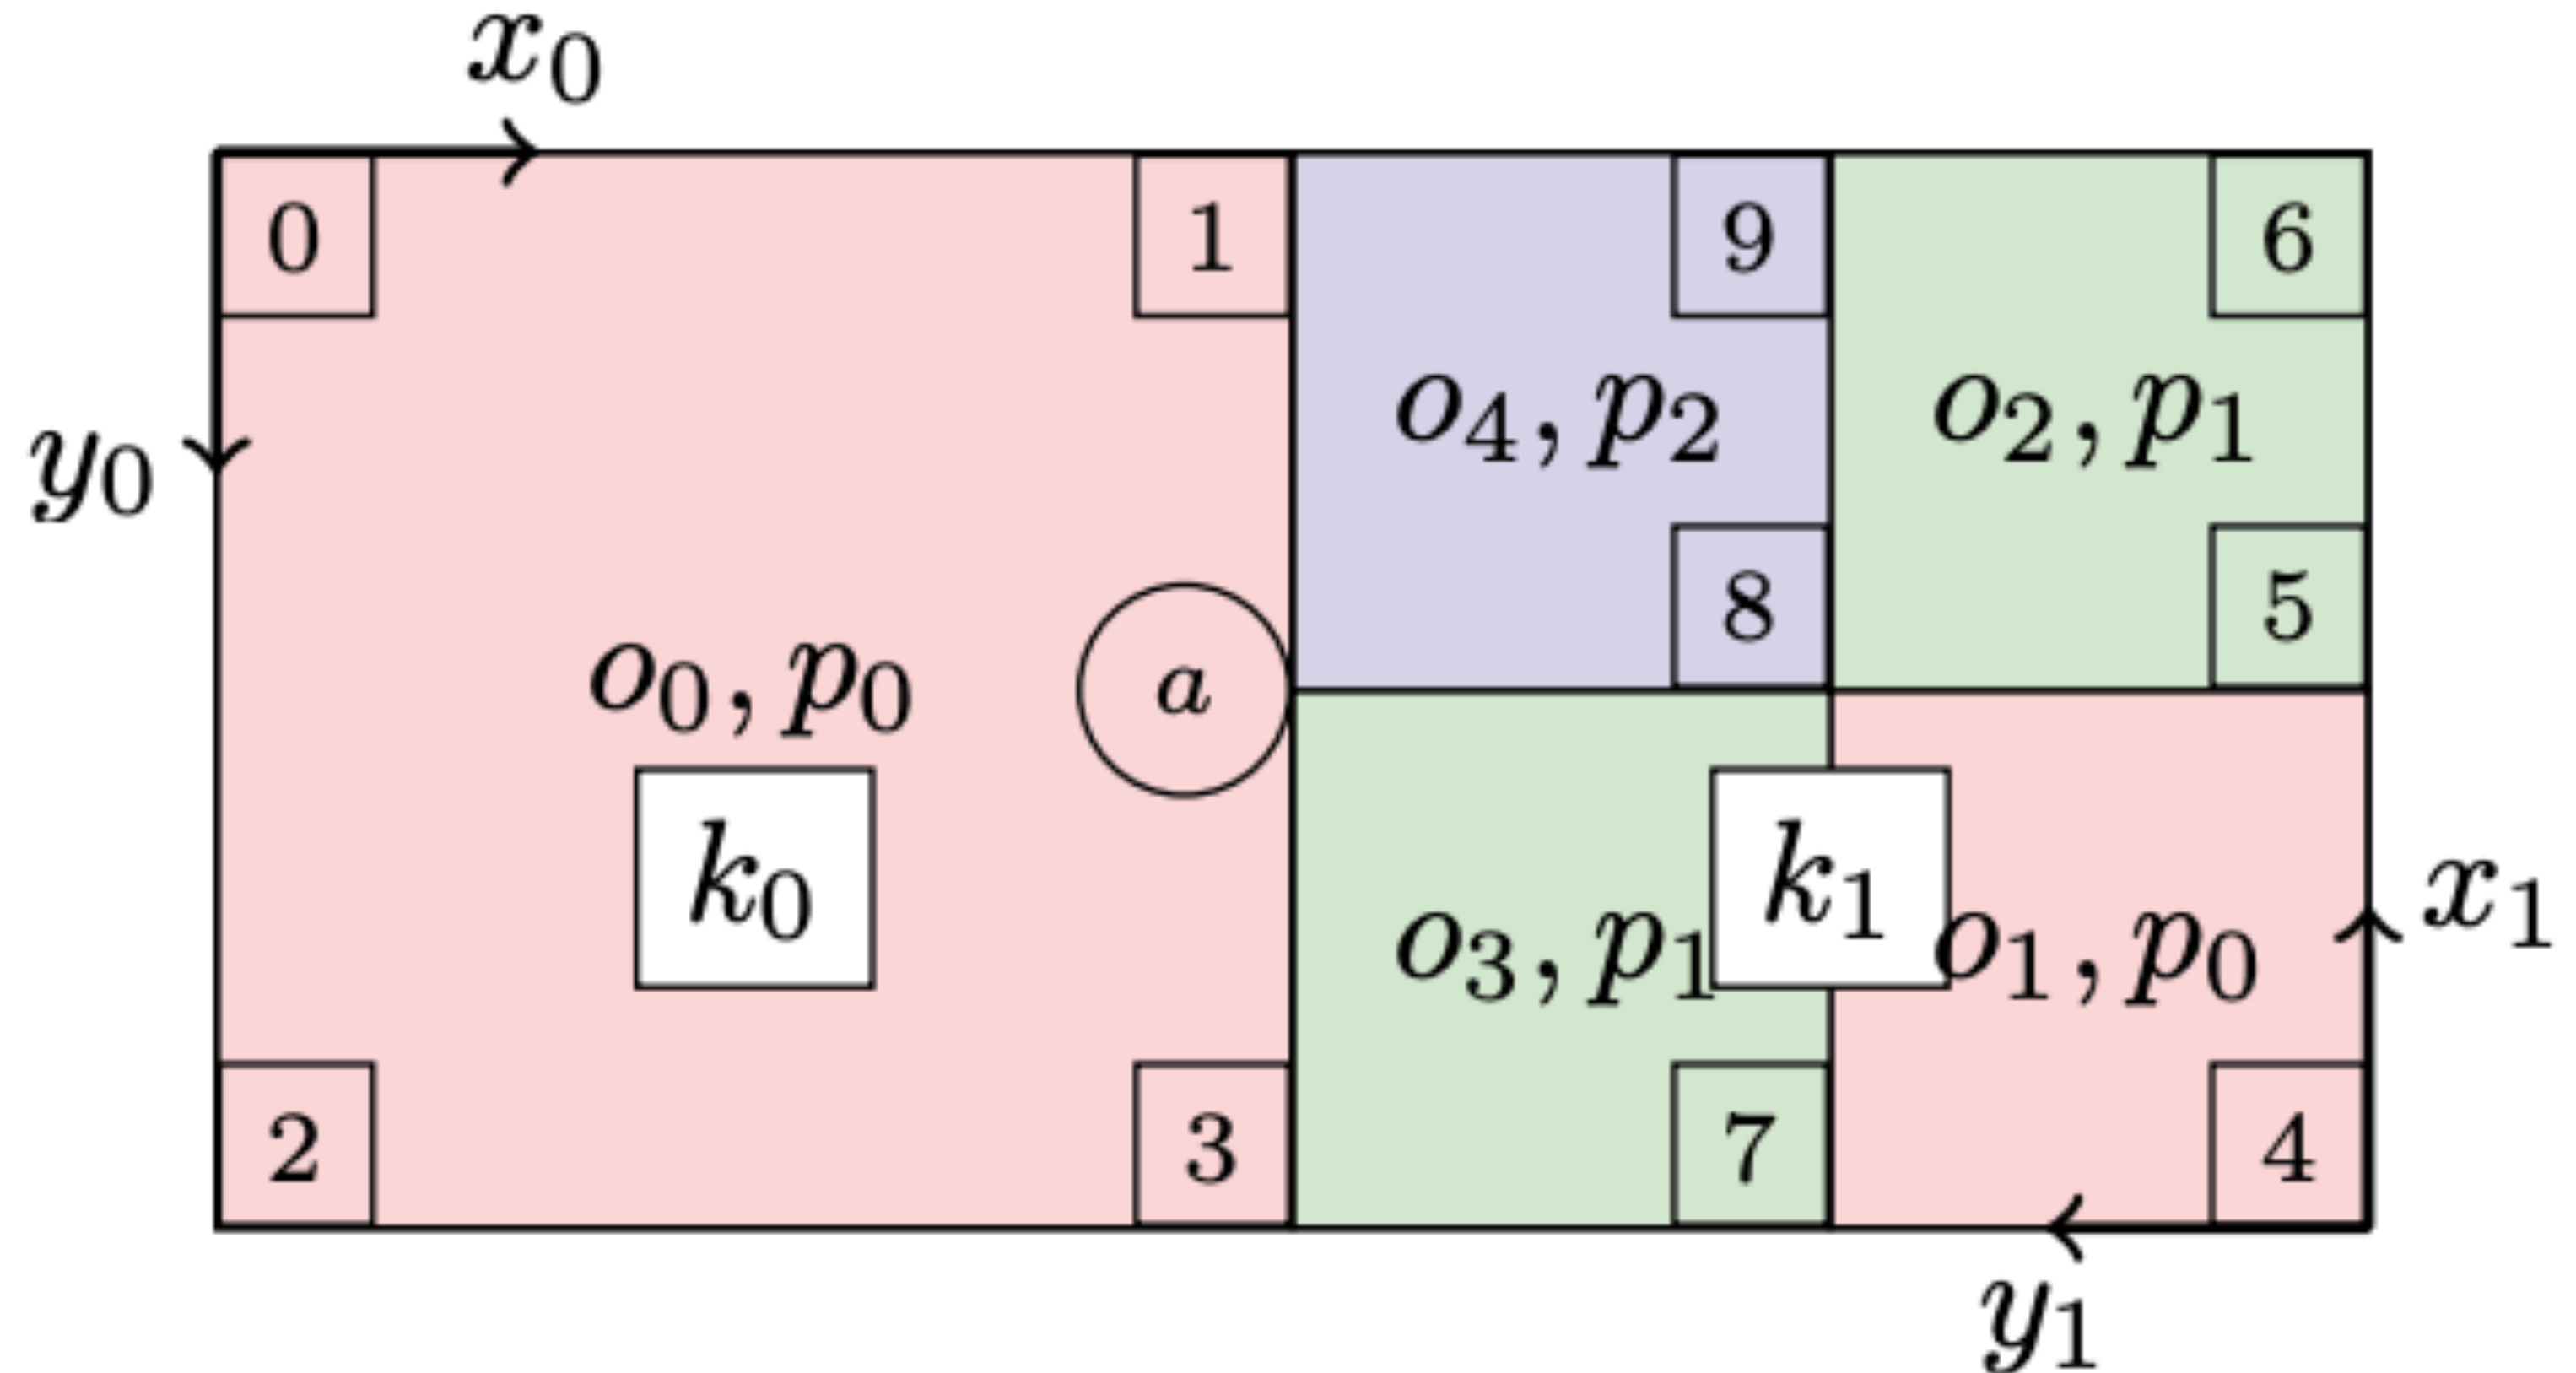
BACKUP

- Development of a **scalable**, **efficient** and **portable** GRMHD code using **GPUs** and **Discontinuous Galerkin** methods.
- Based on the **p4est** library for an efficient, low level forest-of-oct-trees AMR infrastructure.
- GPU offloading handled by the **Kokkos** library, with full support for **HIP** and **CUDA** and experimental support for **Sycle**.
- **Custom** AMR routines and algorithmic flow of the code tailored to target system architectures.
- **Low-level** components of the code (memory allocation, numeric kernels...) redesigned to perform well on new systems.

# Example algorithm: ghost-zone exchange

BACKUP

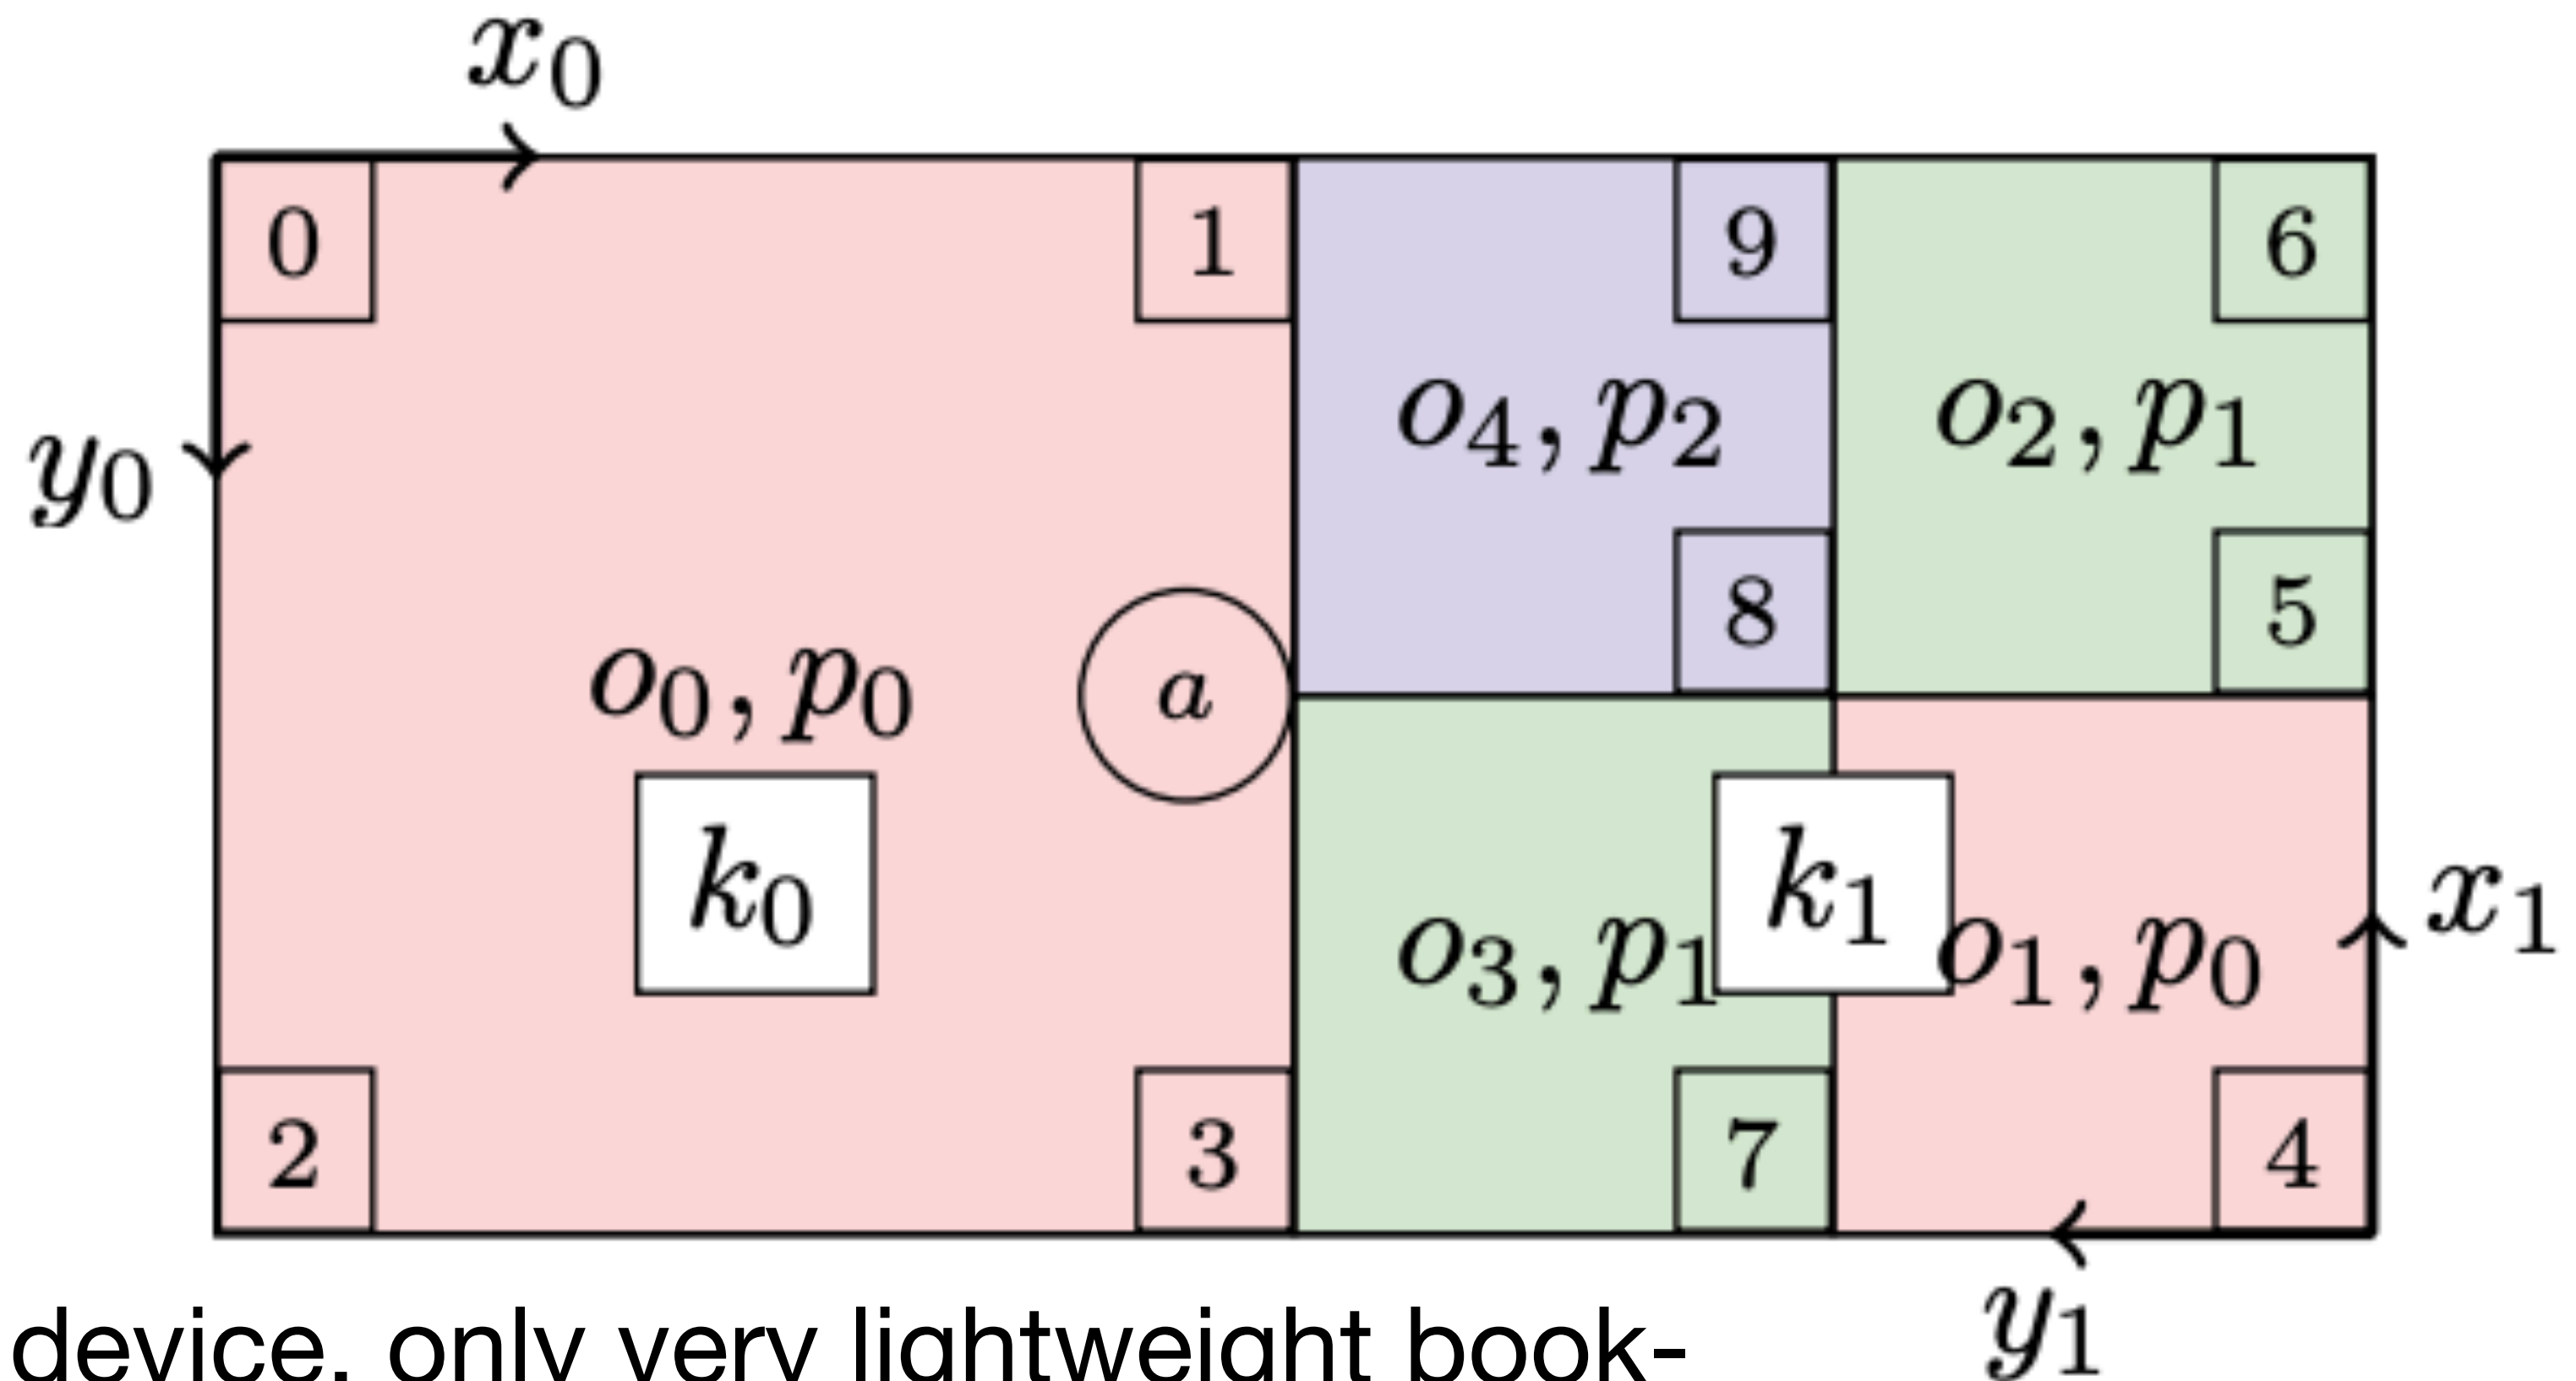
- **1 (Host):** p4est loop over quadrant faces to find neighbors
- **2 (Host-Device):** initiate asynchronous data exchange for halo quads
- **3 (Device):** Apply physical BCs while waiting for data



# Example algorithm: ghost-zone exchange

BACKUP

- **4 (Device):** copy interior simple ghost zones and fill coarse ghost cells from fine data
- **(Host-Device):** Exchange coarse data with filled ghost zones
- **6 (Device):** Fill fine ghost zones from coarse data



Variables **always** sit on device, only very lightweight book-keeping data is exchanged (integers)